

Ada202x Reference Manual – Draft 0.2

S. Tucker Taft, AdaCore

June 7, 2021

Contents

1	Introduction	3
1.1	Language Design Principles	3
1.2	Syntactic Sugar	5
1.3	Relationship between Ada202x, Ada 2022, and ParaSail	5
1.3.1	Ada Features omitted from this subset/variant	5
1.3.2	Features added to this Ada 2022 variant	6
1.4	History of Revisions	6
2	Lexical Elements	7
2.1	Character Set	7
2.2	Delimiters	7
2.3	Identifiers	7
2.4	Literals	8
2.4.1	Integer literals	8
2.4.2	Real Literals	9
2.4.3	Character Literals	9
2.4.4	String Literals	9
2.4.5	Enumeration Literals	10
2.4.6	Null Literals	10
2.5	Comments	10
2.6	Reserved Words	10
3	Types and Objects	11
3.1	Types and Subtypes	11
3.1.1	Record Types and Extensions	12
3.1.2	Syntactic Equivalences	13
3.2	Objects	15
3.3	Object References	16
4	Names and Expressions	17
4.1	Names	17
4.1.1	Component Selection	17
4.2	Expressions	17
4.2.1	Unary and Binary Operators	18
4.2.2	Membership and Null Tests	20
4.2.3	Other Ada202x Operators	20
4.2.4	Aggregates	20
4.2.5	Quantified Expressions	21
4.2.6	Conditional Expressions	22
4.2.7	Map-Reduce Expressions	23

4.2.8	Type Conversion	24
5	Statements	25
5.1	Statement Separators	25
5.2	Assignment Statements	26
5.3	If Statements	26
5.4	Case Statements	27
5.5	Block Statements	28
5.6	Loop Statements	28
5.7	Exit statements	29
6	Operations	30
6.1	Operation Declarations	30
6.2	Operation Types	31
6.3	Operation Definitions	31
6.4	Operation Calls	33
6.5	Return Statements	34
7	Packages	35
7.1	Package Specification	35
7.1.1	Polymorphic Types	37
7.2	Package Body	37
7.3	Package and Type Instantiation	39
8	Containers	40
8.1	Object Indexing and Slicing	42
8.1.1	Syntactic Equivalences	42
8.2	Container Aggregates	43
8.2.1	Syntactic Equivalences	44
8.3	Container Element Iterator	45
8.3.1	Syntactic Equivalences	45
8.4	Container Specifiers	46
9	Aspect Specifications	47
10	synchronized Objects and Parallel Execution	50
10.1	synchronized Types	50
10.1.1	Locked and Queued Operations	51
10.2	Parallel Evaluation	52
11	Ada202x Source Files and Standard Library	53
11.1	Ada202x Source Files	53
11.2	Ada202x Syntax Shorthands	53
11.3	Ada202x Standard Library	54
12	Appendix: Using the Ada202x Interpreter and Virtual Machine	56
12.1	Ada202x Interactive Debugger	57
12.2	Example of using Ada202x Interpreter	57

Chapter 1

Introduction

The *Ada202x* language described in this manual is a subset and variant of the forthcoming Ada2022 language, and is designed with the principle that if you want programmers to write parallel algorithms, you have to immerse them in parallelism, and force them to work harder to make things sequential. In Ada202x, parallelism is everywhere, and threads are treated as resources like virtual memory – a given computation can use 100s of threads in the same way it might use 100s of pages of virtual memory. Ada202x encourages the *divide-and-conquer* approach to parallel computation where threads are each given their own part of the problem to solve, but also supports concurrent access to *synchronized* objects, using both lock-based and lock-free synchronization.

Ada202x also supports aspect specification, such as preconditions and postconditions, and in fact requires them in some cases if they are needed to ensure that a given operation is safe. In particular, all checks that might normally be thought of as run-time checks (if checked by the language at all) are performed at compile time in Ada202x. This includes uninitialized variables, array index out of bounds, null pointers, race conditions, numeric overflow, etc. If an operation would overflow or go outside of an array given certain parameters, then a precondition is required to prevent such parameters from being passed to the operation. Ada202x is designed to support a *formal* approach to software design, with a relatively static model to simplify proving properties about the software, but with an explicit ability to specify run-time polymorphism where it is needed.

Ada202x has four basic concepts – (Generic) Packages, Types, Objects, and Operations. Every type is a private type, a record type, or an operation type. An object is an instance of some type. An operation operates on objects.

The only global variables allowed are those of a synchronized type. Any object to be updated by an operation must be an explicit `[in] out` parameter to the operation, or be identified as a `global var` to the operation.

Ada202x has no pointers, though it has references, optional and expandable objects, and syntactic sugar for indexing (see below for a further discussion of Syntactic Sugar), which together provide a rich set of functionally equivalent capabilities without any hidden aliasing nor any hidden race conditions.

1.1 Language Design Principles

Below are some of the fundamental language design principles we tried to follow while designing Ada202x. Of course, at times we faced a conflict, so at those times tradeoffs had to be made. Although these are expressed as goals, by and large we believe they have been accomplished in the current design.

- The language should be easy to read, and look familiar to a broad swath of existing programmers, from the ranks of programmers in the Algol/Pascal/Ada/Eiffel family, to the programmers in the C/C++/Java/C# family, to the programmers in the ML/Haskell and Lisp/Scheme communities. Readability is to be emphasized over terseness, and where symbols are used, they should be familiar

from existing languages, mathematics, or logic. Although extended character sets are more available these days, most keyboards are still largely limited to the ASCII, or at best, the Latin-1, character set, so the language should not depend on the use of characters that are a chore to type.

Programs are often scanned backward, so ending indicators should be as informative as starting indicators for composite constructs. For example, “end loop” or “end record Stack” rather than simply “end” or “}”.

- Parallelism should be built into the language to the extent that it is more natural to write parallel code than to write explicitly sequential code, and that the resulting programs can easily take advantage of as many cores as are available on the host computer.
- The language should have one primary way to do something rather than two or three nearly equivalent ones. Syntactic sugar (see section below) should be used to provide higher-level constructs, while keeping the core of the language minimal. Nonessential features should be eliminated from the core language, especially those that are error prone or complicate the testing or proof process. User-defined types and language-defined types should have the same capabilities.
- All code should be parameterizable to some extent, since it is arguable that most code would benefit from being parameterized over the precision of the numeric types, the character code of the strings involved, or the element types of the data structures being defined. In other words, any module can be a generic template or equivalent. But the semantics should be defined so that the parameterized modules can be fully compiled prior to being instantiated.
- The language should be inherently safe, in that the compiler should detect all potential race conditions, as well as all potential runtime errors such as the use of uninitialized data, out of bounds indices, overflowing numeric calculations, etc. Given the advances in static analysis, there is no reason that the compiler cannot detect all possible sources of run-time errors.

Programming is about human programmers clearly and correctly communicating with at least two audiences: 1) other human programmers, both current and future, and 2) a very literally-minded machine-based compiler or interpreter. What is needed is *human engineering*, which is the process of adapting a technology to be most useful to humans, by minimizing opportunities for errors, taking advantage of commonly understood principles, using terminology and symbols consistently and in ways that are familiar, and eliminating unnecessary complexity.

Here are some additional somewhat lower level principles followed during the Ada202x design:

- Full generality should be balanced against testability and provability. In particular, though passing functions and types as parameters is clearly useful, it is arguable whether full upward closures and types as true first-class objects (such as the *class* objects in *Smalltalk*), are useful enough to justify the significant testing and proof burdens associated with such constructs. The more disciplined packaging of type and function provided by statically-typed object-oriented programming can match essentially all of the capability provided by upward closures and types as first-class objects, while providing, through behavioral subtyping and other similar principles, a more tractable testing and proof problem.
- Avoid constructs that require fine-grained asynchronous garbage collection if possible. Garbage collectors are notoriously hard to test and prove formally, and are made even more complex when real-time and multi-processor requirements are added. Mark/release strategies, and more generally region-based storage management, as in the *Cyclone* language, suggest possible alternative approaches.
- Mutual exclusion and waiting for a condition to be true should be automatic as part of calling an operation for which it is relevant. This is as opposed to explicit lock/unlock, or explicit wait/signal. Automatic locking and/or waiting simplifies programming and eliminates numerous sources for errors in parallel programs with inter-thread synchronization. The result is also easier to understand and to prove correct.

1.2 Syntactic Sugar

Rather than building in many fundamentally different kinds of types and type constructors, such as enumeration types, array types, fixed-point types, etc., and many different constructs such as concatenation, indexing, and aggregates (e.g. for arrays or other containers), Ada202x uses the notion of *syntactic sugar* to transform these higher-level concepts into the core capabilities of the language. This syntactic-sugar approach allows great flexibility and extensibility, while keeping the core capabilities of the language very simple.

Within the reference manual, *Syntactic Equivalences* sections indicate where syntactic sugar is applied. These sections can be skipped in the initial reading of the reference manual.

1.3 Relationship between Ada202x, Ada 2022, and ParaSail

Describing the language described as a "subset/variant of Ada" is not the whole story. This Ada202x subset/variant is based on the ParaSail parallel programming language, but with syntax drawn from Ada 2022. The ParaSail parallel programming language itself was designed from scratch to support safe and secure parallel programming.

One way to describe this variant of Ada 2022 is with the following four somewhat ironic characteristics:

- *Mutable objects* with value semantics;
- *Stack-based* heap management;
- *Compile-time* exception handling;
- *Race-free* parallel programming.

These characteristics mean that Ada202x remains statically analyzable while providing a flexible, lean, easy to use parallel language.

1.3.1 Ada Features omitted from this subset/variant

Relative to Ada, Ada202x leaves out the following features :

Discriminants – we plan to support a `discriminant_part` with syntactic sugar at some point

Tagged types – any record or private type can be extended – a run-time *type-id* appears only on polymorphic objects;

Tasks and task types – we are considering having a "begin" operator for synchronized types, which would be called by a pico-thread when the object is created, and would be automatically terminated when the scope exits;

Controlled types – we are considering having a "end" operator for synchronized types, which would be called immediately prior to the object being reclaimed;

Entry families – any parameter to a `queued` operation may be used in a `dequeue` condition;

Constrained vs. Unconstrained subtypes/objects – objects can store any value of their subtype – predicates may be used to limit what values are permitted; in particular, arrays can change in length at run-time in this variant;

Interface types – any private type can be used as an interface to be implemented by another type;

Use all type clause – all operations of a type are implicitly visible for calls (and for using as values of an operation type);

Exceptions – support is currently flakey, though preconditions, null values, or multi-threaded exits/returns can substitute;

Access types, allocators, and aliased objects – minimal support at this point; regular types can use a "null" value, and may allow "not null" to be applied to regular types to disallow such use.

1.3.2 Features added to this Ada 2022 variant

Relative to Ada 2022, Ada202x has the following additional features:

Null Values – In this variant, all types have a *null* value, which may be stored in objects not marked as `not null`; these allow objects to grow and shrink – see 3.2;

Instantiation of generic types – a type declared inside a generic package is termed a *generic type* and may be directly instantiated in Ada202x, rather than requiring that the enclosing package first be explicitly instantiated; – see 3.1;

Implicitly parallel semantics – all expression evaluation in Ada202x has implicitly parallel semantics, in that expressions such as $F(X) + G(Y)$ allow evaluation of $F(X)$ and $G(Y)$ in parallel with one another, and statement and loop semantics are designed to simplify automatic parallelization – see 10.2; 5.6;

User-defined literals – any type may use literals, so long as it has a "from_univ" operator for the appropriate Univ_ type – see 2.4, 4.2, and 4.2.3;

Generalized Case Statements – case statements are generalized to support selecting based on the underlying type of a polymorphic object, as well as more generally any sort of set membership – see 5.4;

Private types completed in body – private types are given their full definition in the body of a package rather than in a private part of the package spec.

1.4 History of Revisions

Draft 0.1 Initial revision.

- Based on Draft 0.7 of Sparkel manual

Draft 0.2 Attempt to make it more nearly match reality

Chapter 2

Lexical Elements

2.1 Character Set

Ada202x programs are written using graphic characters from the ISO-10646 (Unicode) character set, as well as horizontal tab, form feed, carriage return, and line feed. A line feed terminates the line.

```
digit ::= 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9
```

```
binary_digit ::= 0 | 1
```

```
hex_digit ::= digit | A..F | a..f
```

```
extended_digit ::= digit | A..Z | a..z
```

2.2 Delimiters

The following single graphic characters are delimiters in Ada202x:

```
( ) { } [ ] , ; . : & | = < > + - * / ' ?
```

The following combinations of graphic characters are delimiters in Ada202x:

```
|| != =? <= >=  
==> ** => [[ ]] << >>  
:= <== <=> <|= += -= *= /= **= <<= >>= |= &=  
.. <.. ..< <..<
```

The following combinations of graphic characters have special significance in Ada202x:

```
and= or= xor=
```

2.3 Identifiers

Identifiers start with a letter, and continue with letters or digits, optionally separated by underscores.

```
identifier ::= letter { [ _ ] ( letter | digit ) }
```


Upper and lower case is significant in identifiers, but two identifiers that differ only in case hide one another. Letters include any graphic character in the ISO-10646 character set that is considered a letter. An all-lower-case identifier must not be the same as an Ada202x reserved word (see 2.6).

Examples:

X, A_B, a123, A123, This_Is_An_Identifier, Xyz_1

NOTE: Upper/lower case is significant in identifiers and reserved words. This eliminates some compatibility issues with new reserved words, since reserved words only conflict with identifiers which are in all lower case, which is a rarity in existing Ada and SPARK code.

2.4 Literals

There are six kinds of literals in Ada202x: integer, real, character, string, enumeration, and null. The syntax for these literals is given below.

```
literal ::=
    integer_literal
  | real_literal
  | character_literal
  | string_literal
  | enumeration_literal
  | null_literal
```

2.4.1 Integer literals

Integer literals are by default decimal. Integers may also be written in binary, hexadecimal, or with an explicit base in the range 2 to 36.

Integer literals are of type Univ_Integer.

```
integer_literal ::=
    decimal_integer_literal
  | based_integer_literal
```

```
decimal_integer_literal ::= decimal_numeral
```

```
based_integer_literal ::= decimal_numeral # extended_numeral #
```

```
decimal_numeral ::= digit { [_] digit }
```

```
hex_numeral ::= hex_digit { [_] hex_digit }
```

```
extended_numeral ::= extended_digit { [_] extended_digit }
```

Examples: 42, 1_000_000, 8#0177#

2.4.2 Real Literals

Real literals are by default decimal, with an optional decimal exponent indicating the power of 10 by which the value is to be multiplied. Reals may also be written with an explicit base in the range 2 to 36, with a decimal exponent indicating the power of the base by which the value is to be multiplied.

Real literals are of type Univ.Real.

```
real_literal ::= decimal_real_literal | based_real_literal

decimal_real_literal ::= decimal_numeral . decimal_numeral [exponent]

based_real_literal ::=
    decimal_numeral # extended_numeral . extended_numeral # [exponent]

exponent ::= (e|E)[+|-] decimal_numeral
```

Examples:

```
3.14159, 0.15, 16#F.FFFF_FFFF_FFFF#e+16
```

2.4.3 Character Literals

Character literals are expressed as a pair of apostrophes bracketing a single `unescaped_character`, being any graphical character of the ISO-10646 character set [*Possible addition:* other than backslash, or a single `escaped_character`, being a backslash followed by an `escapable_character` or a hexadecimal character code].

Character literals are of type Univ.Character.

```
character_literal ::= ' character_specifier '

character_specifier ::=
    unescaped_character
```

Examples:

```
'a', '0'
```

2.4.4 String Literals

String literals are a sequence of graphical characters of the ISO-10646 character set enclosed in double quotes. Two double quotes in a row represent a single double quote within the string. [*Possible addition:* The backslash character may appear only as part of an `escaped_character`.]

String literals are of type Univ.String.

```
string_literal ::= " { character_specifier | "" } "
```

Example:

```
"A simple string literal"
```

[*Possible addition:*

```
"This is a multiline message\n and this is the second line."
```

]

2.4.5 Enumeration Literals

Enumeration literals are expressed with a `#` followed by an identifier or reserved word. In the case of an identifier, the `#` may be omitted in a context where a precondition identifies the subset of enumeration literals which are allowed (as when passed to a `"from_univ"` operator – see 4.2 and 4.2.3).

Enumeration literals are of type `Univ_Enumeration`.

```
enumeration_literal ::= [ # ] identifier | # reserved_word
```

Examples:

```
#case, True, Monday
```

2.4.6 Null Literals

Null literals are specified with the reserved word `null`.

```
null_literal ::= null
```

A `null_literal` represents the null value of the type that is determined by context.

2.5 Comments

Comments in Ada202x start with `--` and continue to the end of the line.

Examples:

```
-- According to the Algol 68 report,  
-- comments are for the enlightenment of the human reader.
```

2.6 Reserved Words

The following words are reserved in Ada202x:

<code>abs</code>	<code>end</code>	<code>new</code>	<code>return</code>
<code>abstract</code>	<code>entry</code>	<code>not</code>	<code>reverse</code>
<code>all</code>	<code>exit</code>	<code>null</code>	<code>separate</code>
<code>and</code>	<code>for</code>	<code>of</code>	<code>some</code>
<code>begin</code>	<code>function</code>	<code>or</code>	<code>then</code>
<code>case</code>	<code>if</code>	<code>package</code>	<code>type</code>
<code>constant</code>	<code>in</code>	<code>parallel</code>	<code>until</code>
<code>continue</code>	<code>is</code>	<code>private</code>	<code>when</code>
<code>declare</code>	<code>limited</code>	<code>procedure</code>	<code>while</code>
<code>else</code>	<code>loop</code>	<code>synchronized</code>	<code>with</code>
<code>elsif</code>	<code>mod</code>	<code>rem</code>	<code>xor</code>

All reserved words in Ada202x are in lower case.

Chapter 3

Types and Objects

In Ada202x, every *object* is an instance of some *type*, and every type is either a *private type*, a *record type*, an *operation type*, or an *instantiation* or *derivation* of one of these. In addition, [NYI] various qualifications and aspect specifications may be applied to the type to produce a particular *subtype* of the type.

3.1 Types and Subtypes

A type is declared as follows:

```
type_declaration ::=
    'type' identifier 'is' ['abstract'] type_definition [ aspect_specification ]

type_definition ::=
    record_type_definition
  | private_type_definition
  | operation_type_specifier
  | type_derivation

private_type_definition ::= ['limited' | 'synchronized'] 'private'

type_derivation ::= 'new' type_specifier_list | type_extension

type_extension ::=
    'new' type_specifier_list 'with' record_type_definition
  | 'new' type_specifier_list 'with' private_type_definition

type_specifier_list ::= type_specifier [ { 'and' named_type_specifier } ]

type_specifier ::=
    named_type_specifier | anon_record_type_specifier | operation_type_specifier

named_type_specifier ::= type_name | type_instantiation

type_instantiation ::= type_name '<' [ generic_actuals ] '>'
```

See 3.1.1 for the syntax of a `record_type_definition` and `anon_record_type_specifier`. See 6.2 for the syntax of an `operation_type_specifier`. See 4.1 for the syntax of a `type_name`.

A subtype is declared as follows:

```
subtype_declaration ::= 'subtype' identifier 'is' subtype_indication
```

```
subtype_indication ::= ['not null'] type_specifier [ aspect_specification ]
```

See chapter 9 for the syntax of an aspect_specification.

A **type_declaration** introduces a new named type. If the declaration occurs within a generic package, the type is called a *generic* type. If the type has the same name as the enclosing package, it is the *primary nested type* of the package, and from outside the package, a name that denotes the package also denotes this type, providing a *short-hand* reference for the type. A **subtype_declaration** introduces a name for a renaming or subtype of an existing (sub)type. A **type_specifier** specifies a type using an **anon_record_type_specifier** (see 3.1.1 below), an **operation_type_specifier** (see 6.2), or by specifying a **type_name** that was declared by a **type_** or **subtype_declaration**, possibly providing generic actual parameters if the specified type is generic. A **subtype_indication** specifies a subtype of the type determined by the **type_specifier**, with possible additional qualifications and aspect specifications.

In a **type_derivation**, the first **type_specifier** identifies the *parent* of the new type, and the new type is *derived* from this parent, and *inherits* operations and components from this type (see Inheritance). Any additional **named_type_specifiers** identify *progenitors* of the new type, and the new type must *implement* the operations of these progenitors; it does not inherit the implementation of any operations nor any components from its progenitors (even if they have components). It inherits code and components only from its parent type.

Two **type_specifiers** identify the *same* type if and only if they are defined by structurally equivalent **anon_record_type_specifiers** (see 3.1.1 below) or **operation_type_specifiers** (see 6.2), or they refer to the same original **type_definition** and specify *equivalent generic_actuals*, if any.

Example:

Given a generic package List providing a List type defined as follows (see 7.1):

```
generic
  type Element_Type is private;
package List is
  type List is private;
  function Create return List;
  function Is_Empty(L : List) return Boolean;
  procedure Append(L : in out List; Elem : Element_Type);
  function Remove_First(L : in out List) return Element_Type;
  function Nth_Element(L : aliased in out List; N : Univ_Integer)
    return aliased Element;
end List;
```

Note that in this example, type List is the *primary nested type* of the package List. A specific kind of list may be declared as follows:

```
type Bool_List is new List < Boolean >;
```

This declares a Bool_List type which represents a list of Booleans.

3.1.1 Record Types and Extensions

Most user-defined types in Ada202x are ultimately record types. A (named) record type or record extension can be defined using a **record_type_definition** within a **type_declaration**, according to the following syntax:

```
record_type_definition ::=
  [ 'limited' | 'synchronized' ] 'record'
  { component_specification ';' }
  'end' 'record' [ identifier ]
```

```

component_specification ::=
    component_mode [ identifier ':' ] subtype_indication [ ':' expression ]

component_mode ::= [ 'ref' ] [ 'const' ]

```

An object of a *limited* type may be assigned a value only as part of its declaration; no subsequent assignments to the object as a whole are permitted (though assignments to individual non-limited components *are* permitted). A type is limited if it has a **ref** component, has a component of a limited type, or has the reserved word **limited** or **synchronized** in its definition.

A **record_type_definition** with a **ref** or limited component is required to be explicitly specified as either **limited** or **synchronized**.

If an identifier is omitted from a **component_specification**, the identifier of the (sub)type specified in the **subtype_indication** is presumed, if it is unique. If this identifier is not unique among the set of identifiers for the components, the component is anonymous.

Rather than declaring a named record type, an *anonymous* record type (also called a *tuple* type) may be specified using an **anon_record_type_specifier**:

```

anon_record_type_specifier ::=
    '(' component_specification { ';' component_specification } ')'

```

An anonymous record type is *limited* if it contains a **ref** or limited component. Two anonymous record types are *structurally equivalent* if they have the same number of components with the same component modes and types. They need not have the same component names or default expressions, if any.

A **const** or **ref const** component is read only. Other components are writable if the enclosing record object is writable. Note that a **ref** component refers to a preexisting *target* object, which necessarily outlives the record object containing the ref.

Implementation note: wrapping a component in a record type should not increase the overall size of objects of the type. Although all types may be extended, normal types do not have *run-time tags*. Only *polymorphic* types have run-time tags (called *type-ids* in Ada202x). See 7.1.1.

3.1.2 Syntactic Equivalences

The declaration of a named record type is equivalent to a derivation from an anonymous record type:

```

type R is record
    A : Integer;
    B : Float;
end record R;
— is equivalent to:
type R is new (A : Integer; B : Float);

```

A type declaration for a private or record type may specify a **discriminant_part** which is equivalent to extending from an anonymous record type that has **const** components:

```

type_declaration ::=
    'type' identifier discriminant_part 'is'
    ['abstract'] type_definition [ aspect_specification ]

```

```

discriminant_part ::= anon_type_specifier

```

```

type T(A : Integer; B : Boolean) is private;
— expands into:
type T is new (const A : Integer; const B : Boolean) with private;

```

There are several additional syntactic constructs provided for various language-provided generic types. These are defined in terms of the core syntax as follows:

```

type_definition ::=
  enumeration_type_definition
  | integer_type_definition | float_type_definition | fixed_type_definition
  | array_type_definition

type_declaration ::= synchronized_type_declaration

enumeration_type_definition ::= '(' identifier {' ',' identifier } ')'

    type E is (A, B, C)
    -- expands into:
    type E is new Enum <[#A, #B, #C]>

    type E is (A, B, C); for E use (1, 2, 4)
    -- expands into:
    type E is new Enum.With_Rep
      <[#A => 1, #B => 2, #C => 4]>

integer_type_definition ::= 'range' expression '..' expression

    type Int is range 1 .. 100
    -- expands into:
    type Int is new Integer <1 .. 100>;

float_type_definition ::= 'digits' expression [ 'range' expression '..' expression ]

    type Flt is digits 5 range -10.0 .. 10.0
    -- expands into:
    type Flt is new Float <Digits => 5, Range => -10.0 .. 10.0>;

fixed_type_definition ::= 'delta' expression [ 'range' expression '..' expression ]

    type Fix is delta 0.01 range -10.0 .. 10.0
    -- expands into:
    type Fix is new Fixed <Delta => 0.01, Range => -10.0 .. 10.0>;

array_type_definition ::=
  'array' '(' index_subtype_definition{' ',' index_subtype_definition } ')' 'of' subtype_specifier

    type Enrollment is array (Course, Semester) of T;
    -- expands into:
    type Enrollment is new Array <T, Indexed_By => (Course; Semester)>>
      -- note use of anon_record_type_specifier here

synchronized_type_declaration ::=
  'synchronized' 'type' identifier 'is'
  { synchronized_operation_declaration }
  'private'
  { synchronized_element_declaration }
  'end' identifier

```

```

synchronized type PT is
  function F return A;
  procedure P (X : Integer);
  entry E (B : Boolean);
private
  D : T;
end PT;
-- expands into:
type PT is synchronized private;
function F (P : in PT) return A;
procedure P (P : in out PT; X : Integer);
procedure E (P : in out PT; B : Boolean);
private
type PT is synchronized record
  D : T;
end record;

```

3.2 Objects

Objects contain data, and may either be variables (declared with 'var'), allowing their data to be changed after initialization, or constants (declared with 'const'), meaning the value of the data of the object cannot be changed between its initial and final reference.

An object is declared using the following syntax:

```

object_declaration ::=
  uninitialized_object_declaration
  | initialized_object_declaration

uninitialized_object_declaration ::=
  identifier ':' ['constant'] ['not null'] subtype_indication [aspect_specification] ';';

initialized_object_declaration ::=
  identifier ':' ['constant'] ['not null'] subtype_indication
    ':= ' expression [aspect_specification] ';';
  identifier ':' 'constant' ':= ' expression [aspect_specification] ';';
  | identifier [':' subtype_indication] 'renames' object_name ';';

var_or_const ::= 'var' | 'const'

```

If an object is declared with a `subtype_indication` without specifying 'not null', its value may be any value of its type that satisfies the requirements of the subtype, as well as the *null* value of its type. An object declared `not null` must be explicitly initialized before use to a non-null value satisfying the requirements of its subtype.

An uninitialized object has the null value initially, but this value cannot be read if the object is declared `not null`. If the uninitialized object is a *constant* without `not null`, this initial null value may be read, but the constant may not then be explicitly assigned a value on the path where the null value is read – the general rule is that two reads of the same constant always return the same value.

Note: as a general rule, if an object is declared `not null`, then its value might nevertheless be null before its first assignment and after its last use, but this null value can never be read.

Examples:


```

BL : not null Bool_List := Create;
T : constant Boolean := True;
Result : T;
Next renames Tree.Left;

```

These declare a variable boolean list, a constant with Boolean value True, a variable Result with implicit initial value of null, and a renaming of Tree.Left as Next.

3.3 Object References

A reference to an existing (*target*) object is declared using the following syntax:

```

object_reference_declaration ::=
  identifier [':' type_specifier ] 'renames' object_name ';'

```

A reference allows the same access as that provided by the object to which it refers.

Examples:

```

const Left renames L.Left_Subtree;
X renames M[I];
Max renames First_Element(A);

```

These create a reference to the Left.Subtree component of L, a reference to the Ith element of M, and a reference to the first element of A, which is a read-write reference only if A is a variable. Note that in the third example, it is assumed that the First.Element function takes an **aliased** parameter and returns an **aliased** result (see 6.1);

Chapter 4

Names and Expressions

4.1 Names

Names denote packages, types, objects, and operations.

```
name ::= package_name | type_name | object_name | operation_name

package_name ::= [ package_name '.' ] identifier

type_name ::= [ package_name '.' ] identifier [ 'attribute_identifier' ]

object_name ::=
    [package_name '.' ] identifier
    | object_indexing_or_slicing
    | operation_call
    | component_selection
```

See Operation Calls (Section 6.4) for the syntax of `operation_name` and `operation_call`. See Object Indexing and Slicing (Section 8.1) for the syntax of `object_indexing_or_slicing`.

4.1.1 Component Selection

If an object is of a record type, or of a type with one or more record extensions, then the components defined by the record type or the record extensions may be named using a `component_selection`. Components are named by naming the enclosing object, then a `'.'`, and then the identifier of the component:

```
component_selection ::= object_name '.' identifier
```

Examples:

```
C.Real_Part , Point.X, List_Node.Next, T.Right_Subtree
```

4.2 Expressions

```
expression ::=
    literal
    | object_name
    | initial_value_specification
```

```

| unary_operator expression
| expression binary_operator expression
| membership_test
| null_test
| quantified_expression
| type_conversion
| type_name
| [ type_name ''' ] bracketed_expression

bracketed_expression ::=
    aggregate
  | conditional_expression
  | map_reduce_expression
  | universal_conversion
  | '(' expression ')'
```

The `null_literal` (the reserved word `null`) evaluates to the *null* value, which is implicitly convertible to any type, and can be used to initialize any object declared to have an `'optional'` type.

Other literals evaluate to a value of a corresponding *universal* type, and are implicitly convertible to any type that has a corresponding `"from_univ"` operator, so long as the value satisfies the precondition of the operator (see 4.2.3).

A `type_name` followed by an apostrophe (`'`) may be used to specify explicitly the result type of a `bracketed_expression` – one of the forms of expression that is enclosed in `()` or `[]`, where the type might not be resolvable without additional context.

A `type_name` by itself is permitted if the type has a `"[.]"` operator defined for itself, and is equivalent to the set of possible values of the type. If the `type_name` identifies a subtype with a non-trivial Predicate, then the set is reduced to those elements of the type that satisfy the Predicate (see Syntactic Equivalences below).

See Aspect Specifications (Chapter 9) for the syntax of `universal_conversion`. See 4.2.7 for the syntax of `initial_value_specification`.

Examples:

```

Y := "This_is_a_string_literal";  — Y must be of a type with a "from_univ" operator
                                — from Univ.String
return null;  — function must have a return type of the form "optional T"
              — indicating it might return "null" rather than a value of type T

Display(Output, Complex'(Real => 1.0, Imaginary => 1.0));
              — Explicitly specify the result type of an aggregate
```

4.2.1 Unary and Binary Operators

The following are the unary operators in Ada202x:

```
"+", "-", "abs", "not"
```

The following are the binary operators in Ada202x:

```

"*)"          -- Exponentiation

"*, "/", "rem", "mod"  -- Multiply, Divide, Remainder, and Modulo operators

"+", "-"      -- Addition and subtraction
```

```

"&"          -- Used to concatenate containers and elements

".. "        -- Interval operator

"<", "<=", "=", -- The usual relational and equality operators
"/=", ">=", ">"

"and", "or", "xor"      -- The basic boolean operators
"and then", "or else"   -- Short-circuit boolean operators

```

The highest precedence operators are the unary operators and the exponentiation ("**") operator. The next lower precedence operators are the multiplication, division, and remainder operators. The next lower precedence operators are the addition, subtraction, and concatenation operators. Next is the interval operator. Next the relational and equality operators. Lowest are the boolean operators.

Addition, subtraction, multiplication, division, and concatenation are left-associative. Exponentiation is right-associative. For other operators, parentheses are required to indicate associativity among operators at the same level of precedence, except that for the boolean operators, a string of uses of the same operator do not require parentheses, and are treated as left-associative.

The binary *compare* operator ("= ?") returns an Ordering value indicating the relation between the two parameters, being **Less**, **Equal**, **Greater**, or **Unordered**. The value **Unordered** is used for types with only a partial ordering. For example, the "= ?" operator for sets would typically return **Equal** if the sets have the same members, **Less** if the left operand is a proper subset of the right, **Greater** if the left operand is a proper superset of the right, and **Unordered** otherwise. All of the other relational operators are defined in terms of "= ?" – only "= ?" is user-definable for a given type.

The evaluation of an expression using a unary or binary operator is in general equivalent to a call on the corresponding operation, meaning that the operands are evaluated in parallel and then the operation is called (see 6.4). The short-circuit boolean operators "and then" and "or else" are implemented in terms of the corresponding *if_expression* (see 4.2.6):

```

A and then B    — equivalent to (if A then B else False)
A or else B     — equivalent to (if A then True else B)

```

Examples of unary and binary operators:

```

"= ?" (S1, S2)    — Compare S1 and S2,
                  — return Less, Equal, Greater, or Unordered
X ** 3            — X cubed
abs (X - Y)       — absolute value of difference
0 .. Length      — The interval 0, 1, .. Length
(A and B) or C    — parentheses required
A or B or C       — parentheses not required
X * Y + U * V     — parentheses not required

```

The relational operators are defined in terms of "= ?" as follows:

```

A = B  ==>  "= ?" (A,B) in [ Equal ]
A != B ==>  "= ?" (A,B) in [ Less , Greater , Unordered ]
A < B  ==>  "= ?" (A,B) in [ Less ]
A <= B ==>  "= ?" (A,B) in [ Less , Equal ]
A > B  ==>  "= ?" (A,B) in [ Greater ]
A >= B ==>  "= ?" (A,B) in [ Greater , Equal ]

```

4.2.2 Membership and Null Tests

A membership test is used to determine whether a value can be converted to a type, satisfies the predicates of a subtype, or is in a particular interval or set. A null test is used to determine whether a value is the null value. The result of a membership test or null test is of type Boolean.

```
membership_test ::=
    expression [ 'not' ] 'in' expression
  | expression [ 'not' ] 'in' type_name

null_test ::= expression 'is' 'null' | expression 'not' 'null'
```

Examples:

```
X in 3..5      — True if X >= 3 and X <= 5
Y not in T'Class — True if Y is not convertible to T'Class
Red in Primary — True if Red satisfies the predicate for Primary
Z not null     — True if Z does not have a null value
```

4.2.3 Other Ada202x Operators

```
"from_univ" -- invoked implicitly to convert from a value of a universal type
"to_univ"   -- invoked using "[[ expression ]]" to convert to a universal type
            -- and used implicitly to convert to a universal type for operations
            -- that take universal-type parameters
"convert"   -- invoked using "type_name ( expression )" to convert between types
"indexing"  -- invoked by "object [ operation_actuals ]" to index into a container
"slicing"   -- invoked by "object [ operation_actuals ]" to select a slice of a container
"index_set" -- invoked by an iterator to iterate over the elements of a container
"[]"        -- invoked by "[]" to create an empty container; invoked implicitly
            -- by "[ key1 => value1, key2 => value2, ... ]" followed by multiple calls
            -- on "|=" to build up a container given the key/value pairs
"[..]"      -- invoked by "[..]" to create a universal set;
            -- invoked implicitly to turn a type name into the set of its values
```

Examples:

```
X := 42;      — Implicit conversion from Univ_Integer using "from_univ" operator
Print( [[X]] ); — Convert back to Univ_Integer for printing using "to_univ" operator
C[Key]        — The element of C associated with given Key using "indexing" operator
A[X..Y]       — The slice of A going from X to Y using "slicing" operator
[]            — An empty container using "[]" operator
```

4.2.4 Aggregates

Aggregates are used for constructing values out of their constituents. There are two kinds of aggregates: the **record_aggregate** for creating an object of a record type from its named components, and the **container_aggregate**, for creating an object of a container type (see 8.2) from a sequence of elements, optionally associated with one or more keys.

The **record_aggregate** is only available for a visible record type. A **container_aggregate** may be used with any type that defines the appropriate operators (see 8.2).

Aggregates have the following form:

```

aggregate ::= record_aggregate | container_aggregate

record_aggregate ::= '(' record_components ')'

record_components ::= [ record_component { ',' record_component } ]

record_component ::=
    [ identifier '=>' ] expression
    | identifier '<==' object_name

```

See 8.2 Container Aggregates for the syntax of a `container_aggregate`.

In a `record_aggregate`, named components (`record_component` with an `identifier` specified) must follow any positional components (those without an `identifier` specified). If the `'<=='` *move* operation is specified, then the value of the component is moved from the named existing object, leaving it null.

Examples:

```

(X => 3.5, Y => 6.2)      — fully named record_aggregate

(Element, Next => null) — mixed positional and named record_aggregate

List := (Item <== Element, Next <== List);
        — move Element to front of linked list

```

4.2.5 Quantified Expressions

Quantified expressions are used to specify a boolean condition that depends on the properties of a set of values.

A quantified expression has the form:

```

quantified_expression ::=
    '(' 'for' all_or_some quantified_iterator '=>' condition ')'

all_or_some ::= 'all' | 'some'

quantified_iterator ::=
    set_iterator | element_iterator | initial_next_while_iterator

```

See Loop Statements (section 5.6) for the syntax of the various iterator forms.

A `quantified_expression` with the reserved word `all` is True if and only if the condition evaluates to True for all of the elements of the sequence produced by the `quantified_iterator`. A `quantified_expression` with the reserved word `some` is True if and only if the condition evaluates to True for at least one of the elements of the sequence produced by the `quantified_iterator`. It is not specified in what order the evaluations of the condition are performed, nor whether they are evaluated in parallel. The condition might not be evaluated for a given element of the sequence if the value for some other element already determines the final result.

Examples:

```

N_Is_Composite := (for some X in 2..N/2 => N rem X = 0);

Y_Is_Max := (for all I in Bounds(A) => A[I] <= Y);

```

4.2.6 Conditional Expressions

Conditional expressions are used to specify a value by conditionally selecting one expression to evaluate among several.

Conditional expressions are of one of the following forms:

```
conditional_expression ::= if_expression | case_expression
```

If Expression

An `if_expression` has the following syntax:

```
if_expression ::=  
    '(' 'if' condition 'then' expression else_part_expression ')'  
  
else_part_expression ::=  
    { 'elsif' condition 'then' expression } 'else' expression
```

All expressions of an `if_expression` must be implicitly convertible to the same type.

To evaluate an `if_expression`, the initial `condition` and those within the associated `else_part_expression` are evaluated in sequence, and the first one that evaluates to `True` determines the expression to be evaluated (the one following the corresponding `'then'`). If all of the conditions evaluate to `False`, the last expression of the associated `else_part_expression` is evaluated to produce the value of the overall `if_expression`.

Examples:

```
Bigger := (if X > Y then X else Y);  
  
return (if Y = 0 then null else X/Y); — return null if would divide by zero
```

Case Expression

Case expressions have the following form:

```
case_expression ::=  
    '(' 'case' case_selector 'is'  
        case_expression_alternative { ';' '  
        case_expression_alternative } [ ';' '  
        case_expression_default ]  
    ')'  
  
case_expression_alternative ::=  
    'when' choice_list '=>' expression  
    | 'when' identifier ':' type_name '=>' expression  
  
case_expression_default ::=  
    'when' 'others' '=>' expression
```

See Case Statements (section 5.4) for the syntax of `case_selector` and `choice_list`.

All expressions following `'=>'` of a `case_expression` must be implicitly convertible to the same type.

The `choice_list` or `type_name` of each `case_expression_alternative` determines a set of values. If there is not a `case_expression_default`, then the sets associated with the `case_expression_alternatives` must cover all possible values of the `case_selector`. The sets associated with the `case_expression_alternatives` must be disjoint with one another, except if there is a `type_name` that identifies a polymorphic type, in which case earlier alternatives take precedence over later polymorphic alternatives.

To evaluate a `case_expression`, the `case_selector` is evaluated. If the value of the `case_selector` is in a set associated with a given `case_expression_alternative`, the corresponding expression is evaluated. If the value is not a member of any set, then the expression of the `case_expression_default` is evaluated.

If a `case_expression_alternative` includes an identifier and a `type_name`, then within the expression, the identifier has the given type, with its value given by a conversion of the `case_selector` to the given type.

Example:

```
return (case Key =? Node.Key is
  when Less => Search(Node.Left , Key);
  when Equal => Node.Value;
  when Greater => Search(Node.Right , Key));
```

4.2.7 Map-Reduce Expressions

Map-reduce expressions are used to specify a value that is produced by combining a set of values, given an initial value and an operation to be performed with each value.

A map-reduce expression has the form:

```
map_reduce_expression ::=
  '[' 'for' map_reduce_iterator
    [ value_filter ] '=>' expression ']' 'Reduce (reducer_operation, initial_value)

map_reduce_iterator ::=
  set_iterator
  | element_iterator

initial_value ::= expression
```

See Loop Statements (section 5.6) for the syntax of `value_filter` and the various iterator forms.

For the evaluation of a `map_reduce_expression`, first the `initial_value` is evaluated and it becomes the *initial* result of the `map_reduce_expression`. Then for each element of the sequence of values produced by the `map_reduce_iterator` that satisfies the `value_filter`, if any, the `expression` is evaluated, and it is combined with the *current* result of the `map_reduce_expression` using the `reducer_operation`, and the result of the evaluation call being the *next* result of the `map_reduce_expression`. After all of the elements of the sequence produced by the iterator have been combined, the last such evaluation determines the *final* result. If there are no elements in the sequence, then the *initial* result is used.

Examples:

```
Sum-Of-Squares := (for X in 1 .. N => <0> + X**2);
```

```
Largest-In-Absolute-Value :=
  (for each E of Arr => Max (<null>, abs E));
```

Note that the language-provided Max operations, when given a null operand, will return the other operand. The same applies to the Min operations.

Syntactic Equivalence

A map-reduce expression is equivalent to a loop that accumulates a result. For example, a map-reduce expression to return the sum of the squares of the odd integers $\leq N$ expands as follows:


```

    return (for X in 1 .. N when X mod 2 = 1 => <0> + X**2);
— expands into:
    Result : Result_Type := 0;

    for X in 1 .. N when X mod 2 = 1 loop
        Result := Result + X**2;
    end loop;
    return Result;

```

4.2.8 Type Conversion

A type conversion can be used to convert an expression from one type to another, by using a syntax like that of an operation call but with the operation identified by the name of the target type:

```
type_conversion ::= type_name '(' expression ')'
```

The expression of a `type_conversion` must be *convertible* to the target type. An expression of a type A is *convertible* to a type B if the type A is *convertible* to type B and the value of the expression after conversion satisfies any value-constraints on B.

Type A is *convertible* to type B if and only if:

- Types A and B are derived from the same original `type_definition` with equivalent actuals if any, but without any type extensions.
- Type B is a polymorphic type (see 7.1.1), and type A is derived from a type equivalent to the root type of B, in the case the root type of B is not a generic type, A implements all of the primitive operations of the root type of B.
- Type A is a polymorphic type, and the run-time type of the expression identifies a type that is convertible to B;
- Type A has a "to_univ" operator and type B has a "from_univ" operator such that the result type of the "to_univ" operator is the parameter type of the "from_univ" operator;
- Type A or type B has a "convert" operator that has a parameter type that matches type A and a result type that matches type B.

Chapter 5

Statements

Statements specify an action to be performed as part of a sequence of statements. A Ada202x statement can either be a simple statement, a compound statement containing other statements as constituents, or a local declaration:

```
statement ::= simple_statement | [ label ] compound_statement | local_declaration

simple_statement ::=
    assignment_statement
  | exit_statement
  | return_statement
  | operation_call

label ::= statement_identifier ':'

statement_identifier ::= identifier

compound_statement ::=
    if_statement | case_statement | loop_statement | block_statement | parallel_block_statement

local_declaration ::= object_declaration | operation_declaration | operation_definition
```

If and only if a `compound_statement` is preceded by a label, then the `statement_identifier` must appear again at the end of the `compound_statement`.

5.1 Statement Separators

Statements are separated with `;`, or with a new line character if the following line is at the same level of indentation. The delimiter `;` may also be used as a statement terminator. The scope of a `local_declaration` occurring immediately within a `statement_sequence` goes from the declaration to the end of the immediately enclosing `statement_list`.

For the execution of a `statement_list` expressions are evaluated and assignments and calls are performed in an order consistent with the order of references to unsynchronized objects (see chapter 10) occurring in the statements.

Examples:

```
A := C(B); D := F(E); U := G(V); W := H(X);
```

```
A : Vector<Integer> := [X, Y];
```

```
parallel do
  Process(A[1]);
and
  Process(A[2]);
end do;
```

The declaration of A is completed before beginning the two separate threads invoking Process on the two elements of A.

5.2 Assignment Statements

An assignment_statement allows for replacing the value of one or more objects with new values.

```
assignment_statement ::=
  object_name ':= ' expression
```

Examples:

```
X := A + B;           — Set X to sum of A and B
```

5.3 If Statements

If statements provide conditional execution based on the value of a boolean expression.

If statements are of the form:

```
if_statement ::=
  'if' condition 'then'
    statement_list
  [ else_part ]
  'end if'

else_part ::=
  'elsif' condition 'then'
    statement_list
  [ else_part ]
  | 'else'
    statement_list
```

```
condition ::= expression -- must be of a boolean type
```

For the execution of an if_statement, the condition is evaluated and if True, then the statement_list of the if_statement is executed. Otherwise, the else_part, if any, is executed.

For the execution of an else_part, if the else_part begins with elsif, then the condition is evaluated and if True, the statement_list following then is executed. Otherwise, the nested else_part, if any, is executed. If the else_part begins with else, then the statement_list following the else is executed.

Example:

```
if This_Were(A_Real_Emergency) then
  You_Would(Be_Instructed , Appropriately);
elsif This_Is(Only_A_Test) then
  Not_To_Worry;
end if;
```

5.4 Case Statements

Case statements allow for the selection of one of multiple statement lists based on the value of an expression.

Case statements are of the form:

```
case_statement ::=
  'case' case_selector 'is'
    case_alternative
    { case_alternative }
    [ case_default ]
  'end' 'case' [ statement_identifier ] [ with_values ]

case_selector ::= expression

case_alternative ::=
  when choice_list '=>' statement_list
  | when identifier ':' type_name '=>' statement_list

choice_list ::= choice { '|' choice }

choice ::= expression [ interval_operator expression ]

interval_operator ::= '..'

case_default ::=
  'when' 'others' '=>' statement_list
```

The `choice_list` or `type_name` of each `case_alternative` determines a set of values. If there is not a `case_default`, then the sets associated with the `case_alternatives` must cover all possible values of the `case_selector`. The sets associated with the `case_alternatives` must be disjoint with one another, except if there is a `type_name` that identifies a polymorphic type, in which case earlier alternatives take precedence over later polymorphic alternatives.

For the execution of a `case_statement`, the `case_selector` is evaluated. If the value of the `case_selector` is in a set associated with a given `case_alternative`, the corresponding `statement_list` is executed. If the value is not a member of any set, then the `statement_list` of the `case_default` is executed.

If a `case_alternative` includes an identifier and a `type_name`, then within the `statement_list`, the identifier has the given type, with its value given by a conversion of the `case_selector` to the given type.

Example:

```
case Lookahead(Input) is
  when 'a'..'z' | 'A'..'Z' =>
    Handle_Alphabetic(Input);
  when '0'..'9' =>
    Handle_Numeric(Input);
  when Control('J') =>
    Handle_End_Of_Line(Input);
  when Control('D') =>
    Handle_End_Of_Input(Input);
  when others =>
    Handle_Others(Input);
end case;
```

5.5 Block Statements

A block statement allows the grouping of a set of statements with local declarations and an optional set of assignments to perform if it completes normally.

A block statement has the following form:

```
block_statement ::=
  [ 'declare'
    declaration_list ]
  'begin'
    statement_list
  'end' [ statement_identifier ]
```

For the execution of a `block_statement`, the `declaration_list` if any is elaborated. Then the `statement_list` is executed.

5.6 Loop Statements

A loop statement allows for the iteration of a `statement_list` over a sequence of objects or values.

Loop statements have the following form:

```
loop_statement ::=
  while_loop | for_loop | indefinite_loop

while_loop ::= 'while' condition loop_body
```

For the execution of a `while_loop` the condition is evaluated. If the condition is evaluates to True then the `statement_list` of the `loop_body` is executed, and if the `statement_list` reaches its end, the process repeats. If the condition evaluates to False, then the `while_loop` is complete.

```
indefinite_loop ::= loop_body
```

An `indefinite_loop` is equivalent to a `while_loop` that has a condition of True.

```
for_loop ::=
  ['parallel [' (' chunk_specification ')']
  'for' iterator [ value_filter ] loop_body

value_filter ::= 'when' condition

loop_body ::=
  'loop'
    statement_list
  'end' 'loop' [ statement_identifier ]

iterator ::=
  set_iterator
  | element_iterator

set_iterator ::=
  identifier [ ':' type_name ] 'in' ['reverse'] expression
```

See 8.3 for the syntax of an `element_iterator`.

The identifier of an iterator declares a *loop variable* which is bound to a particular object or value for each execution of the `statement_list` of the `loop_body`.

Each kind of iterator produces a sequence of values (or objects). If a `value_filter` is present, the sequence is reduced to those values (or objects) that satisfy the `value_filter condition`.

The values in the sequence produced by a `set_iterator` are all of the values of the set, less those that do not satisfy the `value_filter`, if any. See section 8.3 for a description of the sequence of objects, or key-value pairs, produced by an `element_iterator`.

If the `expression` of a `set_iterator` is a `type_name`, it is equivalent to invoking the `'Range'` attribute defined for that (sub)type, to produce the set of all values of the (sub)type (see section 4.2.3).

Examples:

```
parallel
for I in 1..10 loop
  X[I] := I ** 2;
end loop;
```

The above loop initializes a table of squares in parallel.

```
for each S of List_Of_Students(Classroom) when Is_Undergraduate(S) loop
  Print(Report, Name(S));
end loop;
```

The above loop prints the names of the undergraduate students (i.e. those satisfying the `Is_Undergraduate` filter) in the order returned by the `List_Of_Students` function for the given `Classroom`.

5.7 Exit statements

An exit statement may be used to exit a compound statement while terminating any other threads active within the compound statement.

An exit statement has the following form:

```
exit_statement ::=
  'exit' [ statement_identifier ]
  [ 'when' condition ]
```

An exit statement exits the specified `loop_statement` (or in the absence of a `statement_identifier`, the immediately enclosing `loop_statement` of), terminating any other threads active within the identified statement. If the `exit_statement` has a `'when' condition` clause, then the exit is only performed if the condition evaluates to True.

Example:

```
Found : Atomic<Integer> := Create(0);

Search:
parallel
for I in 1 .. 100 loop
  if Matches (A[I], Desired) then
    Set_Value (Found, I);
    exit Search;
  end if;
end loop Search;
```

The above loop searches for an element of an array `A[I]` that matches the `Desired` value, and exists with the first one it finds, after saving the index in an atomic variable.

Chapter 6

Operations

Operations are used to specify an algorithm for computing a value or performing a sequence of actions. There are two kinds of operations – functions (**functions**) and procedures (**procedures**). In addition, a function or a procedure may be an *operator*, as indicated by its designator being an **operator_symbol** (which has the syntax of a **string_literal**). Operators have special meaning to the language, and are invoked using special syntax. Non-operator functions and procedures are invoked using a name followed by parameters (if any) in parentheses. Functions produce one or more results. Operations may update one or more of their variable parameters.

6.1 Operation Declarations

Operations are declared using the following forms:

```
operation_declaration ::=
    function_declaration | procedure_declaration

function_declaration ::=
    'function' designator [ parameter_profile ]
    'return' result_specification

procedure_declaration ::=
    'procedure' designator [ parameter_profile ]

designator ::= identifier | operator_symbol

operator_symbol ::= string_literal

parameter_profile ::=
    '(' parameter_specification { ';' parameter_specification } ')

parameter_specification ::=
    identifier_list ':' [ 'aliased' ] [ parameter_mode ] [ 'not null' ] subtype_indication
    [ ':' expression ]

parameter_mode ::=
    'in' | 'in out' | 'out'
```

```

result_specification ::=
  [ 'aliased' ['constant']] ['not null'] subtype_indication

```

If there is no `parameter_mode`, then the formal is read-only within the body of the operation.

A result indicated as **aliased constant** must be specified via a return statement as a reference to all or part of some 'aliased' parameter (`in`, `in out`, `orout`). A result indicated as **aliased** must be specified via a return statement as a reference to all or part of an **aliased out** or **in out** parameter. *Examples:*

```

function Sin (X : Float) return Float;

function "=?" (Left , Right : Set) return Ordering;

procedure Update (Obj : in out T; New_Info : Info_Type);

function "indexing"(C : aliased Container; Index : Index_Type)
  return aliased Element_Type;

```

6.2 Operation Types

An operation type may be used as a parameter type, to allow operations to be passed as parameters to other operations. Operation-type parameters are considered to be of mode 'in' and do not permit assignment. Operation types are specified with the following syntax:

```

operation_type_specifier ::=
  'access' 'function' [ parameter_profile ] 'return' result_specification
  | 'access' 'procedure' [ parameter_profile ]

```

Two operation types are *structurally equivalent* if they specify the same number of parameters and results, with the same types and modes. Parameter names and defaults are not considered.

Examples:

```

type Trig_Func is access function(Angle : Float) return Float;

type Action_Proc is access procedure(Obj : T);

```

6.3 Operation Definitions

An operation may be defined with a body, with an operation import, or by equivalence to an existing operation.

An operation definition has the following form:

```

operation_definition ::=
  function_definition
  | procedure_definition
  | operation_import
  | operation_equivalence

function_definition ::=
  function_declaration 'is'
    operation_body
  'end' [ 'function' ] designator

procedure_definition ::=

```



```

procedure_declaration 'is'
    operation_body
'end' [ 'procedure' ] designator

operation_body ::=
    [declaration_list
    'begin']
    statement_list

operation_import ::=
    operation_declaration 'with' 'Import' => import_id

import_id ::= string_literal | enumeration_literal

operation_equivalence ::=
    operation_declaration 'renames' operation_name

```

If an operation is declared with a separate `operation_declaration` (typically in a `package_specification`), then the `operation_declaration` part of the `operation_definition` must fully conform to it.

An `operation_import` indicates that the operation is defined externally to the current program, possibly in a different language.

An `operation_equivalence` indicates that the operation is merely a renaming of some existing operation, identified by the `operation_name`. The existing operation must have the same number of parameters and results, of the same modes and with the same types.

Examples:

```

function Sin(X : Float) return Float with Import => "sinf";
    — defined externally

function "+"(Left : Set; Right : Element) return Set renames "&";
    — defined by equivalence

procedure Update(Obj : in outT; New_Info : Info_Type) is
    Obj.Info := New_Info;
end Update;

function Fib (N : Integer) return Integer is
    — Recursive fibonacci but with linear time

    function Fib_Helper(M : Integer)
        return (Prev_Result : Integer; Result : Integer) is
        — Recursive "helper" routine which
        — returns the pair ( Fib(M-1), Fib(M) )
        if M <= 1 then
            — Simple case
            return (Prev_Result => M-1, Result => M);
        else
            — Recursive case
            const Prior_Pair := Fib_Helper(M-1);

            — Compute next fibonacci pair in terms of prior pair
            return with
                (Prev_Result => Prior_Pair.Result ,
                 Result => Prior_Pair.Prev_Result + Prior_Pair.Result);
        end Fib_Helper;
    end Fib;

```

```

        end if;
    end Fib_Helper;

    — Just pass the buck to the recursive helper function
    return Fib_Helper(N).Result;
end function Fib;

```

6.4 Operation Calls

Operation calls are used to invoke an operation, with parameters and/or results.

Operation calls are of the form:

```
operation_call ::= operation_name [ '(' operation_actuals ')' ]
```

```

operation_name ::=
    [ package_name '.' ] operation_designator
    | type_name 'attribute_id'
    | object_name '.' operation_designator
    | object_name 'attribute_id'

```

```
operation_designator ::= operator_symbol | identifier
```

```
operation_actuals ::= operation_actual { ',' operation_actual }
```

```

operation_actual ::=
    [ identifier '=>' ] actual_object
    | [ identifier '=>' ] actual_operation

```

```
actual_object ::= expression
```

```
actual_operation ::= operation_specification | 'null'
```

Unlike other names, an `operation_name` need not identify an operation that is directly visible. Operations declared within packages other than the current package are automatically considered, depending on the form of the `operation_name`:

- If the `operation_name` is of the form `package_name '.' operation_designator` then only operations in the named package are considered.
- If the `operation_name` is of the form `type_name'attribute_id` then only attributes of the named type are considered.
- If the `operation_name` is of the form `object_name'attribute_id` then only attributes of the named object are considered.
- If the `operation_name` is of the form `object_name '.' operation_designator` then the call is equivalent to

```

package_where_object_type_is_declared '.' operation_designator
    '(' object_name ',' operation_actuals ')'

```

- Otherwise (the `operation_name` is a simple `operation_designator`), all operations with the given designator that are operations of any of the parameter types of the call, or of the expected result type of the call, are considered, along with locally declared (non-operator) operations with the given designator. (*Note that all operations of the parameter and results types are automatically visible. In Ada 2012 this would be as though there were a "use all type T" for each parameter or result type of the call. Ada202x does not have a "use [all] type" clause, as it would be redundant.*)

Any named `operation_actuals`, that is, those starting with `"identifier '=>'"`, must follow any positional `operation_actuals`, that is, those without `"identifier '=>'"`.

For the execution of an operation call, the `operation_actuals` are evaluated (in parallel – see 10.2), as are any default expressions associated with non-global parameters for which no actual is provided. After parallel evaluation of the `operation_actuals`, the body of the operation is executed, and then any results are available for use in the enclosing expression or statement.

If the type of one or more of the operation actuals is polymorphic (see 7.1.1), and the operation is an operation of the root type of the polymorphic type, then the actual body invoked depends on the run-time type-id of the actual if the corresponding formal parameter is *not* polymorphic. If multiple operation actuals have this same polymorphic type, and their corresponding formals are also *not* polymorphic, then their run-time type-ids must all be the same (with one exception – the `"=?"` operator always returns `Unordered` when given two polymorphic operands with different type-ids).

Examples:

```
Result := Fib (N => 3);

Graph.Display_Point (X, Y => Sin(X));

A : Sparse_Array := Create(Bounds => 1..N);
```

6.5 Return Statements

A return statement is used to exit the nearest enclosing operation, and if it is a function, specifying the value to return.

A return statement has the following form:

```
return_statement ::=
    'return'
  | 'return' expression
  | 'return' identifier ':' ['constant'] subtype_indication [':=' expression] ['do'
    statement_list
    'end do']
```

If there is no expression, the immediately enclosing operation must be a procedure. If there a single expression, the immediately enclosing operation must be a function. If there is an identifier, this is an extended return statement, where the return object is created, and then may be further initialized in the `statement_list`.

Examples:

```
return Fib(N-1) + Fib(N-2);

return Result : Pair do
    Result.Quotient := Q;
    Result.Remainder := R;
end do;
```

Chapter 7

Packages

Packages define a logically related group of types, operations, data, and, possibly, nested packages. Packages may be generic, parameterized by types, operations, or values.

Every package has a specification that declares external characteristics of its type and objects. If the specification of a package declares any private types or any non-abstract operations, the package must have a body that defines the internal representation of the private types and the algorithms for the operations

7.1 Package Specification

A package is declared by giving its package *specification*. The specification of a package uses the following syntax:

```
package_declaration ::= package_specification ';' | package_instantiation

package_specification ::=
  [ 'generic'
    { generic_formal_parameter ';' } ]
  'package' package_identifier 'is'
  { package_item }
  'end' [ 'package' ] package_identifier

package_identifier ::= { identifier '.' } identifier

package_item ::=
  type_declaration
  | operation_declaration
  | object_declaration
  | package_declaration

package_instantiation ::=
  'package' package_identifier 'is' 'new' package_name '(' generic_actuals ')'
```

Generic formal parameters have the following form:

```
generic_formal_parameter ::= formal_type | formal_object | formal_operation

formal_type ::= type_derivation
```

`formal_object ::= parameter_specification`

`formal_operation ::= 'with' operation_declaration ['is' operation_name]`

An `object_declaration` that occurs immediately within a package specification or package body must be of a synchronized type (see 10.1).

Example (also used in section 3.1):

```
generic
  type Element_Type is new Assignable<>;
package List is
  type List is private;
  function Create return List;
  function Is_Empty(L : List) return Boolean;
  procedure Append(var L : List; Elem : Element_Type);
  function Remove_First(L : in out List) return Element_Type;
  function Nth_Element(L : aliased in out List; N : Univ_Integer) return aliased Element;
end List;
```

This defines the specification of the List package, which defines a type List and operations for creating a list, checking whether it is empty, appending to a list, removing the first element of the list, and getting a reference to the Nth element of the list.

A type may be derived from an existing type, with or without extending the type, and may be defined to *implement* one or more other private types.

When a type T2 is *derived* from a type T1, it inherits *operations* from T1. As part of inheriting an operation from T1, the types of the non-polymorphic (see 7.1.1) parameters and results of the operation are altered by replacing each occurrence of the original type T1 with the new type T2. For example, an operation such as `function Invert(X : T1) return T1` becomes `function Invert(X : T2) return T2`.

An operation inherited from T1 is *abstract* only if the corresponding operation in T1 is abstract, or if the operation has a result which is of a type based on T1 (as does `Invert` in the above example). If the operation inherited from T1 is *not* abstract, then its implicit body is defined to call the operation of T1, with any parameter to this operation that is of the type T1 being passed the parent part of the corresponding parameter to the inherited operation.

An *inherited* operation may be *overridden* by providing a declaration for the operation in the package where the derived type is declared, with the same name and number and types of parameters and results as the inherited operation. An *abstract* inherited operation must be overridden unless the new type is itself specified as **'abstract'**.

Finally, if T1 has any *components*, then if T2 is *derived* from T1, T2 also inherits these components, with any visible components of T1 becoming visible components of T2.

If rather than being derived from T1, the type T2 *implements* T1 (directly or indirectly), and T2 is not itself declared as an abstract type, then T2 is required to declare a corresponding operation for each non-null operation of type T1, but with the change in types of parameters and results from T1 to T2, as described above for inheritance.

If T1 has any *visible* components, then T1 cannot be *implemented* by other types, though T1 may still be *extended*.

Example:

```
generic
  type Skip_Elem_Type is private;
  Initial_Size : Univ_Integer := 8;
package Skip_List is
  type Skip_List is new Lists.List<Element_Type => Skip_Elem_Type>
    and Sets.Set<Elem_Type => Skip_Elem_Type> with private;
```

— The following operations are implicitly declared

```

— due to being inherited from List<Skip_Elem_Type>:

— abstract function Create return Skip_List;
— function Is_Empty(L : Skip_List) return Boolean;
— procedure Append(var L : Skip_List; Elem : Skip_Elem_Type);
— function Remove_First(var L : Skip_List) return optional Skip_Elem_Type;
— function Nth_Element(ref L : Skip_List; N : Univ_Integer)
—   return ref optional Skip_Elem_Type;

function Create return Skip_List;
  — This overrides the abstract inherited operation

... — Here we may override other inherited operations
     — or introduce new operations

function Add(var L : Skip_List; Elem : Skip_Elem_Type) is Append;
  — An operation required by the Set type, defined
  — in terms of one inherited from List.

end Skip_List;

```

7.1.1 Polymorphic Types

If the name of a type is of the form `identifier'Class`, it denotes a *polymorphic* type. A polymorphic type represents the identified type plus any type that extends the type, or that implements all of the identified type's operations, with matching generic actuals. The identified type is called the *root* type for the corresponding polymorphic type.

For example, given the `Skip_List` generic type from the example in ??, and the `Bool_List` type from section 3.1:

```

type Bool_Skip_List is new Skip_Lists.Skip_List<Boolean>;

BL : Bool_List'Class := Bool_Skip_List'Create;

```

The variable `BL` can now hold values of any type that extends or implements the `List` type with `Element_Type` specified as `Boolean`. In this case it is initialized to hold an object of type `Bool_Skip_List`.

An object of a polymorphic type (a *polymorphic object*) includes a *type-id*, a run-time identification of the (non-polymorphic) type of the value it currently contains. The type-id of a polymorphic object may be tested with a membership test (see 4.2.2) or a case statement (see 5.4), and it controls which body is executed in certain operation calls (see 6.4). In the above example, the type-id of `BL` initially identifies the `Bool_Skip_List` type.

7.2 Package Body

A package body defines local types, operations, and objects for a package, as well as the full type for any private type not completed in the private part of the package specification, and a body for each operation declared in the package's specification that requires an implementation.

A package body has the following form:

```

package_body ::=
  'package' 'body' package_identifier 'is'
    { package_body_item }
  'end' [ 'package' ] package_identifier ','

```

```

package_body_item ::=
  package_item
| operation_body
| package_body

```

Example:

```

package body List is
  type List_Node is record
    var Elem : Element_Type;
    var Next : optional List_Node;
  end record List_Node;

  type List is record
    var Head : optional List_Node;
  end record;

  function Create return List is
    return (Head => null);
  end Create;

  function Is_Empty(L : List) return Boolean is
    return L.Head is null;
  end Is_Empty;

  procedure Append(var L : List; Elem : Element_Type) is
    for X => L.Head loop
      if X is null then
        -- Found the end, add new component here
        X := (Elem => Elem, Next => null);
      else
        -- Iterate with next node
        continue loop with X => X.Next;
      end if;
    end loop;
  end Append;

  function Remove_First(var L : List) return Result : optional Element_Type is
    if L.Head is null then
      -- List is empty, nothing to return
      return null;
    else
      -- Save first element and then delete node from list
      Result := L.Head.Elem;
      L.Head := L.Head.Next;
      return; -- Result already assigned
    end if;
  end Remove_First;

  function Nth_Element(ref L : List; N : Univ_Integer)
  return ref optional Element is
    for (X => L.Head; I := 1) loop
      if X is null then
        -- reached end of list
        return null;
      elsif I = N then

```

```

        -- reached Nth element
        return X.Elem;
    else
        -- continue with next node of list
        continue loop with (X => X.Next, I => I+1);
    end if;
end loop;
end Nth_Element;

end List;

```

The above defines the body for the package List whose specification is given in 7.1. Type List_Node is a type used for the implementation of the exported type List. The full type declaration is provided for the type List declared in the specification of package Lists. Following that are the bodies for the operations exported by the Lists package.

7.3 Package and Type Instantiation

Generic packages and types within them are instantiated by providing actuals to correspond to the generic formals. If an actual is not provided for a given formal, then the formal must have a default specified in its declaration, and that default is used.

The actual parameters used when instantiating a generic package or type to produce a (non-generic) package or type have the following form:

```

generic_actuals ::= [ generic_actual { ',' generic_actual } ]

generic_actual ::=
    [ identifier '=>' ] actual_type
  | [ identifier '=>' ] actual_operation
  | [ identifier '=>' ] actual_object

actual_type ::= subtype_indication

```

Any generic actuals with a specified identifier must follow any actuals without a specified identifier. The identifier given preceding '=>' in a generic_actual must correspond to the identifier of a formal parameter of the corresponding kind.

Chapter 8

Containers

A container is a type that defines an "indexing" operator, an "index_set" operator, a container aggregate operator "`[]`", a combining assignment operator "`|=`", and, optionally, a "slicing" operator. It will also typically define a Length or Count function, other operations for creating containers with particular capacities, for iterating over the containers, etc.

The *index type* of a container type is determined by the type of the second parameter of the "indexing" operator, and the *value type* of a container type is determined by the type of the result of the "indexing" operator.

The *index-set type* of a container type is the result type of the "`index_set`" operator, and must be either a set or interval over the index type.

Examples:

```
generic
  type Key_Type is new Hashable<>;
  type Element_Type is private;
package Map is
  type Map is private;
  function "[]" return Map;
  procedure "|=" (var M : Map; Key : Key_Type; Elem : Element_Type);
  function "indexing" (ref M : Map; Key : Key_Type)
    return ref optional Element_Type;
  function "index_set" (M : Map) return Set<Key_Type>;
end Map;
```

The Map package defines a container type Map with Key_Type as the index type and Element_Type as the value type. The package includes a parameterless container aggregate operator "`[]`" which produces an empty map, a combining operator "`|=`" which adds a new Key => Elem pair to the map, "indexing" which returns a reference to the element of M identified by the Key (or null if none), and "`index_set`" which returns the set of Keys with non-null associated elements in the map.

```
generic
  type Element_Type is new Hashable<>;
package Set is
  type Set is private;
  function "[]" () return Set;      — Empty set
  procedure "|=" (var S : Set; Elem : Element_Type);
  function Count (S : Set) return Univ_Integer;
  function "in" (Elem : Element_Type; S : Set) return Boolean;
  function "indexing" (ref S : Set;
    Index : Univ_Integer)
    return ref Elem
```

```

    with Pre => Index in 1 .. Count(S);
    function "index_set"(S : Set) return Interval<Univ_Integer>;
    function "=?"(Left , Right : Set) return Ordering;
end Set;

```

The Set package defines a container type Set with the Element_Type as the value type and Univ_Integer as the index type. The package includes a parameterless container aggregate operator "[]" which produces an empty set, a combining operator "|=" which adds a new element to the set, an "in" operator which tests whether a given element is in the set, an "indexing" operator which returns the *n*-th element of the set, and an "index_set" operator which returns the interval of indices defined for the set (i.e. 1..Count(S)). The compare operator ("=?" – see 4.2.1) is provided for comparing sets for equality and subset/superset relationships.

```

generic
    type Component_Type is private;
    type Indexed_By is new Countable<>;
package Array is
    type Array is private;
    type Bounds_Type is Interval<Indexed_By>;
    function Bounds(A : Array) return Bounds_Type;

    function "[]"
        (Index_Set : Bounds_Type; Values : Map<Bounds_Type, Component_Type>)
        return Result : Array
        with Post => Bounds(Result) = Index_Set;

    function "indexing"(ref A : Array;
        Index : Indexed_By)
        return ref Component_Type;
    function "index_set"(A : Array)
        return Result : Bounds_Type
        with Pre => Index in Bounds(A),
            Post => Result = Bounds(A);

    function "slicing"(ref A : Array;
        Slice : Bounds_Type)
        return ref Result : Array
        with Pre => Slice <= Bounds(A),
            Post => Bounds(Result) = Slice;

    procedure "|="(var A : Array;
        Index : Indexed_By;
        Value : Component_Type)
        with Pre => Index in Bounds(A);
    procedure "|="(var A : Array;
        Slice : Bounds_Type;
        Value : Component_Type)
        with Pre => Slice <= Bounds(A);
end Array;

```

The Array package defines a container type Array with Component_Type as the value type and Indexed_By as the index type. The index-set type is Bounds_Type. The package includes a container aggregate operator "[]" which creates an array object with the given overall Index_Set and the given mapping of indices to values. It also defines an "indexing" operator which returns a reference to the component of A with the given Index, a "slicing" operator which returns a reference to a slice of A with the given subset of the Bounds, plus combining operators "|=" which can be used to specify a new value for a single component or all components of a slice of the array A.

8.1 Object Indexing and Slicing

Object indexing is used to invoke the "indexing" operator to obtain a reference to an element of a container object. Object slicing is used to invoke the "slicing" operator to obtain a reference to a subset of the elements of a container object.

Object indexing and slicing use the following syntax. The form with '`[..]`' is only for slicing.

```
object_indexing_or_slicing ::=
  object_name '[' operation_actuals ']'
  | object_name '[..]'
```

If the form with '`[..]`' is used, or one or more of the `operation_actuals` are sets or intervals, then the construct is interpreted as an invocation of the "slicing" operator. Otherwise, it is interpreted as an invocation of the "indexing" operator. The `object_name` denotes the container object being indexed or sliced.

When interpreted as an invocation of the "slicing" operator, the construct is equivalent to:

```
"slicing" '(' object_name, operation_actuals ')'
```

or, for the form using '`[..]`':

```
"slicing" '(' object_name ')'
```

When interpreted as an invocation of the "indexing" operator, the construct is equivalent to:

```
"indexing" '(' object_name, operation_actuals ')'
```

The implementation of an "indexing" operator must ensure that, given two invocations of the same "indexing" operator, if the actuals differ between the two invocations, then the results refer to different elements of the container object. Similarly, the implementation of a "slicing" operator must ensure that, given two invocations of the same "slicing" operator, if at least one of the actuals share no values between the two invocations, then the results share no elements.

If an implementation of the "indexing" operator and an implementation of the "slicing" operator for the same container type have types for corresponding parameters that are the same or differ only in that the one for the "slicing" operator is an interval or set of the one for the "indexing" operator, then the two operators are said to *correspond*. Given invocations of corresponding "indexing" and "slicing" operators, the implementation of the operators must ensure that if at least one pair of corresponding parameters share no values, then the results share no elements of the container object. A slice defined using '`[..]`' is presumed to refer to *all* elements of the container.

Examples:

```
Table[Key] += 1;      — bump up Table entry associated with Key

A[1..3] <=> A[4..6]; — swap halves of 6-element array
Qsort(V[..]);         — Pass a slice representing all of V to Qsort
```

8.1.1 Syntactic Equivalences

For compatibility with existing SPARK code, parentheses '`(' ... ')'`' may be used instead of brackets '`[...]'`' for indexing or slicing:

```
M ( I, J )
— is equivalent to
M [ I, J ]

A ( 1 .. 10 )
— is equivalent to
A [ 1 .. 10 ]
```

8.2 Container Aggregates

A container aggregate is used to create an object of a container type, with a specified set of elements, optionally associated with explicit indices.

```
container_aggregate ::=
    empty_container_aggregate
  | universal_container_aggregate
  | positional_container_aggregate
  | named_container_aggregate
  | iterator_container_aggregate

empty_container_aggregate ::= '[]'

universal_container_aggregate ::= '[..]'

positional_container_aggregate ::=
    '[' positional_container_element { ',', positional_container_element } ']'

positional_container_element ::= expression | default_container_element

default_container_element ::= 'others' '=>' expression

named_container_aggregate ::=
    '[' named_container_element { ',', named_container_element } ']'

named_container_element ::=
    choice_list '=>' expression
  | default_container_element

iterator_container_aggregate ::=
    '[' 'for' iterator [ value_filter ] [ direction ] [ ',', index_expr ]
    '=>' expression ']'

index_expr ::= expression
```

An `empty_container_aggregate` is only permitted if the container type has a parameterless container aggregate operator `[]`.

A `universal_container_aggregate` is only permitted if the container type has a universal set operator `[..]`.

The `choice_list` in a `named_container_element` must be a set of values of the index type of the container. The expression in a `container_element` must be of the value type of the container.

If present in a `container_aggregate`, a `default_container_element` must come last. A `default_container_element` is only permitted when the `container_aggregate` is being assigned to an existing container object, or the index-set type of the container has a universal set operator `[..]`.

In an `iterator_container_aggregate`, the iterator must not be an `initial_value_iterator`, and if it is an `initial_next_while_iterator`, it must have a `while_or_until` condition.

The evaluation of a `container_aggregate` is defined in terms of a call on a container aggregate operator `[]` or `[..]`, optionally followed by a series of calls on the combining move operation `<|=` (for `positional_container_aggregates`) or the `"var_indexing"` operator (for `named_container_aggregates`).

For the evaluation of an `empty_container_aggregate`, the parameterless container aggregate operator `[]` is called. For the evaluation of a `universal_container_aggregate`, the parameterless universal container aggregate operator `[..]` is called.

For the evaluation of a `positional_container_aggregate` or a `named_container_aggregate`:

- if there is a container aggregate operator "`[]`" which takes an index set and a mapping of index subsets to values, this is called with the index set a union of the indices defined for the aggregate, and the mapping based on the container elements specified in the `container_aggregate`. The `default_container_element` is treated as equivalent to the set of indices it represents.
- if there is only a parameterless container aggregate operator "`[]`" then it is called to create an empty container; the combining operator "`<|="`" is then called for each `positional_container_element` in the aggregate, while the "`var_indexing`" operator is called for each `named_container_element`, with a `choice_list` of more than one choice resulting in multiple calls.

If there is a `default_container_element`, it is equivalent to a `container_element` with a `choice_list` that covers all indices of the overall container not covered by earlier `container_elements`.

For the evaluation of an `iterator_container_aggregate`, the expression is evaluated once for each element of the sequence of values produced by the iterator, with the loop variable of the iterator bound to that element. If an explicit `index_expr` is present, it is provided as the index to the "`var_indexing`" operator. Otherwise, the loop variable is implicitly provided as the index, unless a `direction` or `value_filter` is specified or no "`var_indexing`" operator is available, in which case the "`<|="`" operator is used to build up the container value without any explicit index. If there is a `value_filter`, the expression is evaluated only for elements of the sequence that satisfy the `value_filter`.

Examples:

```
[ 1, 2, 3, 4, 5 ]           — positional container aggregate
[ 1..5 => 1, others => 0 ]   — named with default
[ Red => 0x1, Green => 0x10, Blue => 0x100 ]
                             — all named
[ for I in 0..10 when I mod 2 = 0 => I ** 2 ] — table of even squares
[ for I in 1..N, Key[I] => Value[I] ] — mapping given key/value vectors
```

8.2.1 Syntactic Equivalences

For compatibility with existing SPARK code, parentheses '`(...)`' may be used instead of brackets '`[...]`' for container aggregates.

A container aggregate is expanded into a series of calls on operators of the container type.

```
[A, B, C]
— expands into:
var Agg : Container_Type := []
Agg <|= A; Agg <|= B; Agg <|= C

[K1 => V1, K2 => V2, K3 => V3]
— expands into:
var Agg : Container_Type := []
"var_indexing"(Agg, K1) := V1
"var_indexing"(Agg, K2) := V2
"var_indexing"(Agg, K3) := V3
```

A use of `others =>` expands into the set of indices of the existing container object or the universal set operator "`[..]`" not already covered by earlier choices.

An `iterator_container_aggregate` is expanded into a loop using either the "`var_indexing`" or "`<|="`" operator to add elements to the container.

```
[for I in 1..10 => I * 2]
— if there is an "var_indexing" operator available, expands into:
var Agg : Container_Type := []
```

```

    for I in 1..10 loop
        Agg[I] := I * 2;
    end loop
— if only a "<|=" operator is available , expands into:
    var Agg : Container_Type := []
    for I in 1..10 forward loop
        — NOTE: "forward" is used by default (since the set "1..10" supports it)
        Agg <|= I * 2
    end loop

    [ for each V of C when V > 0 => V ]
— expands into:
    var Agg : Container_Type := []
    for each V of C when V > 0 [forward] loop
        — NOTE: "forward" is used only if container supports it.
        Agg <|= V
    end loop

    [ for X => First then X.Next while X not null => X.Data ]
— expands into:
    var Agg : Container_Type := []
    for X => First then X.Next while X not null loop
        Agg <|= X.Data
    end loop

```

8.3 Container Element Iterator

An element iterator may be used to iterate over the elements of a container.

An element iterator has the following form:

```

element_iterator ::=
    identifier [ ':' type_name ] 'of' expression
    | '[' identifier '>' identifier ']' 'of' expression

```

An `element_iterator` is equivalent to an iterator over the index set of the container identified by the `expression`. In the first form of the `element_iterator`, in each iteration the `identifier` denotes the element of the container with the given index. In the second form of the `element_iterator`, the first identifier has the value of the index itself, and the second identifier denotes the element at the given index in the container. The `identifier` denoting each element of the container is a variable if and only if the container identified by the `expression` is a variable.

Example:

```

for each [ Key => Value ] of Table loop
    — Iterate over key/value pairs of table
    Display(Output, Key, Value);
end loop;

```

8.3.1 Syntactic Equivalences

An element iterator is equivalent to an iterator over the "index_set" of the container, as follows:

```

for each [ Key => Value ] of Table loop
    ...
end loop

```

```

— expands into:
  for Key in "index_set"(Table) loop
    ref Value => "indexing"(Table, Key) — i.e. Table[Key]
    ...
  end loop

```

An element iterator with only a single identifier is equivalent to the [Key => Value] form with an anonymous Key.

```

  for each Elem of Container loop ...
— expands into:
  for each [Anon_Key => Elem] of Container loop ...

```

8.4 Container Specifiers

There are various operations in Ada202x for moving rather than copying objects and components of objects, such as the "<==" and the "<|=" operations (see section 5.2). These can be used to reduce the amount of copying that is performed, which can be important when dealing with containers whose elements are themselves large objects. In some cases, we may build up a large object, with the intent of moving it into a container. In this case, there is some advantage to indicating, when the object is declared, that it is specifically intended to be moved into a particular container or other object when complete. This will cause its storage to be allocated in the same region as that of the specified container or other existing object.

The container or object whose region is to be used may be indicated when declaring an object or a parameter, using a `container_specifier`, whose syntax is as follows:

```

container_specifier ::= 'for' object_name

```

If a formal parameter of an operation has a `container_specifier`, then the `_specifier's` `object_name` must identify a `var` parameter or a result of the same operation.

A parameter with a `container_specifier` is set to null as a side-effect of the call. This occurs even if the formal parameter is not specified as a `var` parameter. The actual parameter is allowed to be a constant so long as it is the last use of the constant. If the actual parameter is an expression, then the region of the named object becomes a target for the evaluation of the expression. Note that the original value of the actual parameter may be preserved by parenthesizing the actual parameter, presuming the mode of the formal parameter is not `var`. This ensures that a copy of the actual parameter is made at the point of call.

Examples:

```

— Compute the intersection of Left and Right
— and put result back in Left.
var Result : Set for Left := []; — Result in same region as Left
for Elem in Right loop
  if Elem in Left then
    Result |= Elem; — Add Elem to intersection
  end if;
end loop;
Left <== Result; — Move result to be new value for Left.

procedure Move_Into_Set(var Left : Set; Right : Element_Type for Left);
— Value of Right moved into Left,
— leaving Right null.

function Destructive_Union(Left, Right : Set for Result) return Result : Set;
— Left and Right are union'ed to
— form the Result, with Left and Right
— ending up null after the call.

```

Chapter 9

Aspect Specifications

Declarations may be annotated using `aspect_specifications` to specify a precondition of an operation, a postcondition of an operation, a predicate for a subtype or an object, or an invariant for a type.

Aspect specifications have the following form:

```
aspect_specification ::=
  'with' aspect_mark [ '=>' expression ] {',',
    aspect_mark [ '=>' expression ] }
```

Preconditions, postconditions, predicates, and invariants specified by `aspect_specifications` are checked by the Ada202x compiler, and it will complain if it cannot prove that the associated conditions evaluate to True on any possible execution of the program.

```
universal_conversion ::= '[' expression '']'
```

An expression of the form `'[[' expression ']]'` may be used to convert an expression to a universal type, generally for use in an `aspect_specification` for a precondition or a postcondition. The type of the expression must have a `"to_univ"` operator; the type of the `universal_conversion` is the result type of this operator.

Examples:

```
function Sqrt(X : Float) return Float
with Pre => X >= 0.0,
    Post => Sqrt'Result >= 0.0;

type Age is new Integers.Range<0..200>;
subtype Minor is Age with Predicate => Minor < 18;
subtype Senior is Age with Predicate => Senior >= 50;
```

These `aspect_specifications` define predicates on two different subtypes of the Age type.

```
generic
  Modulus : Univ_Integer with Predicate => Modulus >= 2;
package Modular.Types is
  type Mod is private;
  function "from_univ"(Univ : Univ_Integer)
    return Modular
    with Pre => Univ in 0 ..< Modulus;

  function "to_univ"(Val : Modular) return Result : Univ_Integer
    with Post => Result in 0 ..< Modulus;
```



```

function "+"(Left , Right : Modular) return Result : Modular
  with Post => [[Result]] = ( [[Left]] + [[Right]] ) mod Modulus;
  ...
end Modular_Types;

```

The precondition on "from_univ" indicates the range of integer literals that may be used with a modular type with the given modulus. The postcondition on "to_univ" indicates the range of values returned on conversion back to Univ_Integer. The postcondition on "+" expresses the semantics of the Modular "+" operator in terms of the language-defined operations on Univ_Integer.

Here is a longer example:

```

generic
  type Component is private;
  type Size_Type is new Countable<>;
package Stacks is
  type Stack is private;
  function Max_Stack_Size(S : Stack) return Size_Type;
  function Count(S : Stack) return Size_Type;

  function Create(Max : Size_Type) return Stack
    with Pre => Max > 0,
        Post => Max_Stack_Size(Create'Result) = Max
        and Count(Create'Result) = 0;

  procedure Push
    (var S : Stack;
     X : Component)
    with Pre => Count(S) < Max_Stack_Size(S),
        Post => Count(S) = Count(S)'Old + 1;

  function Top(ref S : Stack) return ref Component
    with Pre => Count(S) > 0;

  procedure Pop(var S : Stack)
    with Pre => Count(S) > 0,
        Post => Count(S) = Count(S)'Old - 1;
end Stack;

package body Stacks is
  type Stack is record
    const Max_Len : Size_Type;
    Cur_Len : Size_Type
    with Predicate => Cur_Len in 0..Max_Len;
    Data : Array<optional Component, Indexed_By => Size_Type>
    with Predicate => Length(Data) = Max_Len;
  end record
  with Type_Invariant => (for all I in 1..Cur_Len => Data[I] not null);
    — invariant needed for Top()

  function Max_Stack_Size(S : Stack) return Size_Type is
    return S.Max_Len;
  end function Max_Stack_Size;

  function Count(S : Stack) return Size_Type is
    return S.Cur_Len;
  end function Count;

```

```

function Create(Max : Size_Type) return Stack is
    return (Max_Len => Max, Cur_Len => 0, Data => [1.. Max_Len => null]);
end Create;

procedure Push
    (var S : Stack;
     X : Component) is
    S.Cur_Len += 1;
    S.Data[S.Cur_Len] := X;
end Push;

function Top(ref S : Stack) return ref Component is
    return S.Data[S.Cur_Len];
end function Top;

function Pop(var S : Stack) is
    S.Cur_Len -= 1;
end Pop;
end Stacks;

```

Chapter 10

synchronized Objects and Parallel Execution

Expression evaluation in Ada202x proceeds in parallel (see 6.4), as do statements separated by '||' (see 5.1), and the iterations of a parallel loop (see 5.6 and ??). The Ada202x implementation ensures that this parallelism does not introduce *race conditions*, situations where a single object is manipulated concurrently by two distinct threads without sufficient synchronization. A program that the implementation determines might result in a race condition is illegal.

Objects in Ada202x are either *synchronized* or *unsynchronized*, according to whether their type is or is not a *synchronized* type. *synchronized* objects allow concurrent operations by multiple threads by using appropriate hardware or software synchronization. An *unsynchronized* object allows concurrent operations only on non-overlapping parts of the object.

10.1 synchronized Types

A type is *synchronized* if it has the reserved word **synchronized** in its definition, or if it is derived from a *synchronized* type.

Example:

```
generic
  type Item_Type is new Machine_Integer<>;
package Lock_Free is
  type Atomic is synchronized private;
  function Create(Initial_Value : Item_Type) return Atomic;
  function Test_And_Set(var X : Atomic) return Item_Type;
    — If X = 0 then set to 1; return old value of X
  function Compare_And_Swap(var X : Atomic;
    Old_Val, New_Val : Item_Type) return Item_Type;
    — If X = Old_Val then set to New_Val; return old value of X
end Lock_Free;
...
var X : Lock_Free.Atomic<Int_32> := Create(0);
var TAS_Result : Int_32 := -1;
var CAS_Result : Int_32 := -1;
block
  TAS_Result := Test_And_Set(X);
  ||
  CAS_Result := Compare_And_Swap(X, 0, 2);
end block;
```

- Now either $TAS_Result = 0$, $CAS_Result = 1$, and X is 1,
- or $TAS_Result = 2$, $CAS_Result = 0$ and X is 2.

This is an example of a package which defines a synchronized atomic type whose objects can hold a single `Machine_Integer`, and can support concurrent invocations by multiple threads of `Test_And_Set` and `Compare_And_Swap` operations. The implementation of this type would presumably use hardware synchronization.

10.1.1 Locked and Queued Operations

The operations of a synchronized type may include the reserved word `locked` or `queued` for parameters of the type. If a synchronized type has any operations that have such parameters, then it is a *locking* type; otherwise it is *lock-free*. Any object of a locking type includes an implicit *lock* component.

If an operation has a parameter that is marked `locked`, then upon call, a lock is acquired on that parameter. If it is specified as a `var` parameter, then an exclusive read-write lock is acquired; if it is not specified as a `var` parameter then a sharable read-only lock is acquired. Once the lock is acquired, the operation is performed, and then the caller is allowed to proceed.

If an operation has a parameter that is marked `queued`, then the body of the operation must specify a *dequeue* condition. A *dequeue_condition* has the following form:

```
dequeue_condition ::= 'queued' while_or_until condition 'then'
```

A dequeue condition is *satisfied* if the condition evaluates to True and the reserved word `until` appears, or if the condition evaluates to False and the reserved word `while` appears.

Upon call of an operation with a `queued` parameter, a read-write lock is acquired, the dequeue condition of the operation is checked, and if satisfied, the operation is performed, and then the caller is allowed to proceed. If the dequeue condition is not satisfied, then the caller is added to a queue of callers waiting to perform a queued operation on the given parameter.

Within an operation of a synchronized type, given a parameter that is marked `locked` or `queued`, the components of that parameter may be manipulated knowing that an appropriate lock is held on that object. If there is a synchronized type parameter that is *not* marked `locked` or `queued`, then there is no lock on that parameter, and only synchronized components of such a parameter may be manipulated directly.

If upon completing a locked or queued operation on a given object, there are other callers waiting to perform queued operations, then before releasing the lock, these callers are checked to see whether the dequeue condition for one of them is now satisfied. If so, the lock is transferred to that caller and it performs its operation. If there are no callers whose dequeue conditions are satisfied, then the lock is released, allowing other callers not yet queued to contend for the lock.

If an operation declared in the visible part of a package performs a call on a queued operation internally, but does not have a queued parameter, then the operation as a whole must be marked with the reserved word `'queued'` prior to the reserved word `'function'` or `'procedure'` (see 6.1). This indicates that an indefinite delay within the operation might occur, while waiting for the dequeue condition associated with some call to be satisfied. Such operations must not be called from within a locked operation, as they could cause a lock to be held indefinitely. On the other hand, operations with a parameter explicitly marked as `'queued'` may be called while already holding a lock on that parameter, but the dequeue condition must already be satisfied at the point of call.

Example:

```
generic
  type Element_Type is private;
package Queue is
  type Queue is synchronized private;
  function Create return Queue;
  procedure Append(locked var Q : Queue; Elem : Element_Type);
  function First(locked Q : Queue) return optional Element_Type;
```

```

    — Returns null if queue is empty.
function Remove_First(queued var Q : Queue) return Element_Type;
    — Queued until the queue has at least one element
end Queue;
...
var Q : Queue<Int32> := Create();
var A : Int32 := 0;
var B : Int32 := 0;
block
    Append(Q, 1); Append(Q, 2);
    ||
    A := Remove_First(Q);
    ||
    B := Remove_First(Q);
end block;
— At this point, either A = 1 and B = 2
— or A = 2 and B = 1.

```

In this example, we use a locking type Queue and use locked and queued operations from three separate threads to concurrently add elements to the queue and remove them, without danger of unsynchronized simultaneous access to the underlying queuing data structures.

Note: operations of a Ada202x synchronized type with a queued parameter are similar to Ada's synchronized entries, with the dequeue condition being analogous to the entry barrier. However, a dequeue condition in Ada202x may depend on the value of any parameter to the queued operation, whereas in an Ada synchronized entry, the barrier may not depend on any parameter, though it may depend on an entry family index, if any.

10.2 Parallel Evaluation

Two expressions that are parameters to an operation call (see 6.4) or a binary operator (see 4.2.1) are evaluated in parallel in Ada202x, as are the expressions that appear on the right hand side of an assignment and those within the **object_name** of the left hand side (see 5.2). In addition, the separate **statement_threads** of a **statement_thread_group** (see 5.1) are performed in parallel. Finally, the iterations of a parallel loop (see 5.6) are performed in parallel.

Two **object_names** that can be part of expressions or statements that are evaluated in parallel must not denote overlapping parts of a single unsynchronized object, if at least one of the names is the left-hand side of an assignment or the actual parameter for a **var** parameter of an operation call. Distinctly named components of an object are non-overlapping. Elements of a container associated with distinct indices are non-overlapping (see 8.1).

Examples:

```

function Bump(var A : Int) return Int;
X := 3 || X := 5           — illegal
X := 3 || Y := X           — illegal
A := X || B := X           — legal
A[I] := 2 || A[J] := 3     — illegal if I can equal J
Bump(X) + X                — illegal
X := Bump(X)               — legal

```

Chapter 11

Ada202x Source Files and Standard Library

11.1 Ada202x Source Files

Each Ada202x source file is made up of a sequence of standalone package or operation definitions. `With_clauses` can be used to specify which other packages or operations are visible when defining a given standalone package or operation.

```
source_file ::=
    { compilation_unit }

compilation_unit ::= context_clause standalone_program_unit

context_clause ::= { with_clause [ use_clause ] }

with_clause ::= 'with' program_unit_name { ',' program_unit_name } ';'

use_clause ::= 'use' package_name ';'

standalone_program_unit ::=
    package_declaration | package_body | operation_definition
```

A `with_clause` may be used to control which other standalone program units are visible within a given `standalone_program_unit`. A `use_clause` makes the declarations within a (non-generic) package specification directly visible. For a generic package, a `use_clause` makes the generic types within the package specification directly visible.

Note: There is no "use type" clause in Ada202x. See the description of "operation_name" resolution in 6.4.

11.2 Ada202x Syntax Shorthands

Ada202x syntax is not as rigid as implied by the BNF given in this reference manual. In particular, semicolons at the end of a statement or declaration may be omitted, and "begin" may be omitted between declarations and statements.

For example, the following is a legal use of these shorthands:

```

function Fib (N : Integer) return Integer is
  — Recursive fibonacci but with linear time

  function Fib_Helper(M : Integer)
    return (Prev_Result : Integer; Result : Integer) is
    — Recursive "helper" routine which
    — returns the pair ( Fib(M-1), Fib(M) )
    if M <= 1 then
      — Simple case
      return (Prev_Result => M-1, Result => M)
    else
      — Recursive case
      const Prior_Pair := Fib_Helper(M-1)

      — Compute next fibonacci pair in terms of prior pair
      return with
        (Prev_Result => Prior_Pair.Result ,
         Result => Prior_Pair.Prev_Result + Prior_Pair.Result)
    end if
  end Fib_Helper — This is optional

  — Just pass the buck to the recursive helper function
  return Fib_Helper(N).Result

```

11.3 Ada202x Standard Library

Ada202x includes a number of language-provided packages in the Ada202x Standard Library, with names of the form `A2X.Core.*` and `A2X.Containers.*`.

`A2X.Core` itself is a non-generic package, and provides the predefined types of the language:

Univ_Integer arbitrary length integers

Univ_Real ratio of two `Univ_Integers`, with plus/minus zero and plus/minus infinity

Univ_Character 31-bit ISO-10646 (Unicode) characters

Univ_String vector of `Univ_Characters`

Boolean an enumeration type with two values `False` and `True`

Ordering an enumeration type with four values `Less`, `Equal`, `Greater`, and `Unordered`

The other language-provided types include:

type Any is abstract limited private all types implement `Any` implicitly

type Assignable is abstract private provides `:=`, `<==`, and `<=>` operations; all non-limited types implement `Assignable` implicitly

type Comparable is abstract limited private provides `=?` operator

type Hashable is abstract private implements `Assignable` and `Comparable`; provides `Hash` operation

type Countable is abstract private implements `Hashable`; provides `+` and `-` operators to add or subtract `Univ_Integers` to progress through the values of the type

type Imageable is abstract private implements Hashable; provides **To_String** and **From_String** functions to convert the value to and from a **Univ_String**.

type Enum **<Vector<Univ_Enumeration>>** implements Countable and Imageable; used to define a new enumeration type given the vector of literals; see Syntactic Equivalences subsection of 3.1

type Integer **<Interval<Univ_Integer>>** implements Countable and Imageable; provides the usual operators

type Float<Univ_Integer> implements Hashable and Imageable; provides the usual operators

type Array<Element_Type is Assignable<>; Indexed_By is Countable<>> a fixed-size array of **Element_Type**, indexed by a specified countable type

Chapter 12

Appendix: Using the Ada202x Interpreter and Virtual Machine

The Ada202x interpreter and virtual machine is invoked with the following command:

```
% a2xi [ -debug on ] [ -servers nnn ] [ -listing on/off ] file1.a2x file2.a2x ...  
      [ -command func arg1 arg2 ... ]
```

The `a2xi` command invokes the interpreter on the specified Ada202x source files. If `-command` is specified, then after translating the Ada202x source to ParaSail Virtual Machine (PSVM) instructions, it executes the specified Ada202x `func` with the given arguments, if any. If `-command` is not specified, a prompt is presented once processing is successful, allowing the user to enter a command to the virtual machine in the form "`func arg1 arg2 ...`".

At the interactive prompt, the `quit` command may be used to exit the virtual machine. The "`debug on`" command may be used to turn on virtual machine debugging. The "`debug off`" command may be used to turn it back off. The "`servers <number>`" command may be used to specify the total number of server threads to use for work stealing, henceforth.

The number of servers can also be specified on the command line via "`-servers <number>`". The default is six. That is, there is effectively an implicit "`-servers 6`" on the command line if no explicit "`-servers nnn`" is specified. At the interactive prompt, "`servers nnn`" is only effective to increase the number of servers, because once a server is activated there is no mechanism to deactivate it.

To simplify the use of the interpreter, we have provided a c-shell script which has fewer options, but which also automatically incorporates the new interactive Ada202x debugger console (`lib/debugger_console.psi`), which will be invoked if an assertion or pre/postcondition fails during execution:

```
% ../bin/ada_interp.csh file1.a2x file2.a2x ... [-command func arg1 arg2 ...]
```

If you do *not* use `ada_interp.csh`, but instead use `a2xi` directly, the first source file to be interpreted should *always* be the Ada202x standard library, with name "`../lib/aaa.psi`". This file contains the definitions for the standard modules such as `Univ_Integer`, `Set`, `Vector`, etc. When using `ada_interp.csh`, this file, along with the Ada202x debugger sources (which are actually in ParaSail), are included automatically.

If the Ada202x interpreter detects an error during parsing, semantic analysis, or PSVM code generation, it gives error messages on the standard error output stream, as well as creating a file "`errors.err`," which may be viewed using the "`vim`" text editor giving it the "`-q`" flag. When "`errors.err`" is viewed using "`vim -q`", it will automatically position the text window at the line in a source file with a compilation error. The `vim` commands `:cn` and `:cp` may be used to go to the next and the previous error.

In addition to producing error messages, the Ada202x interpreter also produces listing files, with names of the form "`file1.a2x.lst`". These include a line-numbered listing of the source file, an *unparsing* of each

top-level compilation unit, and the ParaSail Virtual Machine instructions generated for each operation. These listing files are produced whether or not an error is detected. The ParaSail Virtual Machine instructions are only produced if the Ada202x compiler's semantic analysis of the source code succeeds. By default, a listing is not produced when there is a `"-command"` specified. This default can be overridden with the `"-listing on/off"` option on the `a2xi` command line. Giving `"-listing off"` will turn off listings in all cases. Giving `"-listing on"` will produce a listing in any case. When using the `ada_interp.csh` script, you will need to give the `-w` flag to have a listing produced (see `ada_interp.csh -h` for a full list of flags for this script).

Note that the error recovery of this Ada202x interpreter is not perfect, so it may be necessary to fix the first few errors and then rerun `a2xi` or `ada_interp.csh` to avoid being misled by cascading errors.

12.1 Ada202x Interactive Debugger

There is now an interactive debugger which is loaded automatically if you use the `ada_interp.csh` script to run the Ada202x interpreter. The interactive debugger is invoked whenever the interpreter encounters an assertion or a pre/postcondition that fails at run-time. It is also invoked when the interpreter hits some other sort of run-time failure.

The debugger console is itself written in Ada202x, and is in `lib/debugger_console.a2x`. Feel free to take a look at how it works. When it is invoked, giving the command `help` (or `h`) will list the available commands:

```
>> Debugger command: help
```

Debugger commands available:

```
quit|q|exit : terminate the program and exit
continue|c|con|cont : continue execution
up|u : go to next outer stack frame, if any
down|d : go to next inner stack frame, if any
list|l [+|-|<line>] : list lines from source file
params|par : print values of parameter for current frame
locals|loc : print value of locals of current frame
help|h : show this message
```

The `up` and `down` commands walk up and down the stack of whatever is the current server thread running at the time of the failure. This can be confusing because in the presence of `concurrent` loops or calls on locking or queuing operations of `concurrent` objects, the stack frame reached by `up` might *not* be the logically immediately enclosing stack frame. It *does* illustrate how *work stealing* works, where a single (heavyweight) server thread executes bits and pieces (pico-threads) of a given program. In a future release, we plan to change `up` and `down` so they walk up and down the logical hierarchy of stack frames, and allow explicit switching between different logical threads of control. An adventurous user might try experimenting with `lib/debugger_console.a2x` to see if they can enhance the debugger in this or some other way. Please send us the results of your experiments!

12.2 Example of using Ada202x Interpreter

Examples of using the Ada202x interpreter:

```
% ada_interp.csh qsort.a2x
Ada202x Interpreter and Virtual Machine Revision: 9.0
Copyright (C) 2011-2021, AdaCore, New York NY, USA
This program is provided "as is" with no warranty.
```

```

Parsing <install-dir>lib/aaa.a2i
Parsing /<install-dir>/lib/reflection.psi
Parsing <install-dir>/reflection.a2x
Parsing <install-dir>/psvm_debugging.a2x
Parsing <install-dir>/lib/debugger_console.a2x
Parsing qsort.a2x
---- Beginning semantic analysis ----
Starting up thread servers
  162 trees in library.
Done with First pass.
Done with Second pass.
Installing Debugging Console!
Done with Pre codegen pass.
Done with Code gen.
Filling in cur-inst-param info in op tables.
Evaluating global constants.
Finishing type descriptors.

Command to execute: Test_Sort 10
Before sort, Vec =
  70, 43, 1, 92, 65, 26, 40, 98, 48, 67
After sort, Vec =
  1, 26, 40, 43, 48, 65, 67, 70, 92, 98
After 2nd sort, Vec2 =
  1, 26, 40, 43, 48, 65, 67, 70, 92, 98

Command to execute: quit
Shutting down thread servers

Stg_Rgn Statistics:
New allocations by owner:      3312  = 55%
Re-allocations by owner:      1892  = 31%
Total allocations by owner:    5204  = 87%

New allocations by non-owner:   317   = 5%
Re-allocations by non-owner:   440   = 7%
Total allocations by non-owner: 757   = 12%

Total allocations:              5961

Threading Statistics:
Num_Initial_Thread_Servers : 1 + 1
Num_Dynamically_Allocated_Thread_Servers : 4
Max_Waiting_Shared_Threads (on all servers' queues): 1
Average waiting shared threads: 0.00
Max_Waiting_Unshared_Threads (on any one server's queue): 1
Average waiting unshared threads: 0.00
Max_Active (threads): 6
Average active threads: 2.20
Max_Active_Masters : 6
Max_Subthreads_Per_Master : 3

```

```

Max_Waiting_For_Subthreads : 3
Num_Thread_Steals : 570 out of 1 + 570 (U+S) thread initiations = 99%

% ada_interp.csh qsort.a2x -command Test_Sort 10
Before sort, Vec =
  70, 43, 1, 92, 65, 26, 40, 98, 48, 67
After sort, Vec =
  1, 26, 40, 43, 48, 65, 67, 70, 92, 98
After 2nd sort, Vec2 =
  1, 26, 40, 43, 48, 65, 67, 70, 92, 98

% ada_interp.csh error_test.a2x
Ada202x Interpreter and Virtual Machine Revision: 9.0
Copyright (C) 2011-2021, AdaCore, New York NY, USA
...
Parsing /Users/stt/_parasail/lib/aaa.a2i
...
Parsing error_test.a2x
error_test.a2x:7:14: Error: Use ":=" rather than "=" in Ada202x
error_test.a2x:9:8: Error: Use "/=" rather than "!="
error_test.a2x:11:7: Error: Use "elsif" rather than "elseif"
error_test.a2x:11:12: Error: Use "=" rather than "=="
error_test.a2x:34:30: Error: Use "=" rather than "=="
error_test.a2x:48:4: Error: Start label Misspelling does not match end label Misspelling
error_test.a2x:57:14: Error: Use ":=" rather than "=" in Ada202x
error_test.a2x:66:19: Error: Use "/=" rather than "!="
error_test.a2x:69:5: Error: Use "end if" rather than "endif"
error_test.a2x:79:15: Error: Syntax Error
error_test.a2x:80:1: Error: Should be "end function"
  11 syntax errors found in error_test.a2x

  11 total syntax errors found
---- All done ----

% vim -q
[interactive correction of errors identified in errors.err]

```